

PREFACE TO THE SECOND EDITION

in the years since the first edition of this book appeared, Artificial Intelligence (AI) has grown from scale laboratory science into a technological and industrial success. We now possess an arsenal of techniques for creating computer programs that control manufacturing processes, diagnose computer faults and human diseases, design computers, do insurance underwriting, play grandmaster-level chess and so on. Basic research in AI has expanded enormously during this period. For the student extracting theoretical and practical knowledge from such a large body of scientific knowledge is a daunting task. The goal of the first edition of this book was to provide a readable introduction to the problems and techniques of AI. In this edition, we have tried to achieve the same goal for the expanded field that AI has, and; become. In particular, we have tried to present both the theoretical foundations of AI and an indication of the ways that current techniques can be used in application programs,

As a result of this effort, the book has grown. It is probably no longer possible to cover everything in a single semester. Because of this, we have structured the book so that an instructor can choose from a variety of paths through the chapters. The book is divided into three parts:

Part I. Problems and Search.

Part II. Knowledge Representation.

Part III. Advanced Topics,

Part I introduces AI by examining the nature of the difficult problems that AI seeks to solve. It then develops the theory and practice of heuristic search, providing a detailed comparison of standard search methods, including best-first search, hill climbing, simulated annealing, and constraint satisfaction.

The last thirty years of AI have demonstrated that intelligence requires more than the ability to reason. It also requires *glean* deal of knowledge about the world. So Part II explores a variety of methods for encoding knowledge in computer systems. These methods include predicate logic, production rules, semantic networks, frames, and scripts. There are also chapters on both symbolic and numeric techniques for reasoning under uncertainty. In addition, we present some very specific frameworks in which particular commitments to a set of representational primitives are made.

Parts I and II should be covered in any basic course in AI. They provide the foundation for the advanced topics and applications that are presented in Part III. While the chapters in Parts I and II should be covered in order since they build on each other, the chapters in Part III are the most independent and can be covered in almost any combination, depending on the goals of a particular course. The topics that are covered include: game playing, planning, understanding, natural language processing (which depends on the undemanding chapter), parallel and distributed AI (which depends on planning and natural language), teaming, connectionist models, common sense expert systems, and perception and action.

To use this book effectively, students should have some background in both computer science and mathematics. As computer background, they should have experience programming and they should feel comfortable with the material in an undergraduate data structures course. They should be familiar with the use of recursion as a program control structure. And they should be able to do simple analyses of the time complexity of algorithms. As mathematical background, students should have the equivalent of an undergraduate course in logic, including predicate logic with quantifiers, and the basic notion of a decision procedure.

This book contains, spread throughout it, many references to the AI research literature. These references are important for two reasons. First, they make it possible for the student to pursue topics in greater depth than is possible within the space restrictions of this book. This is the common reason for including references in a survey text. The second reason that these references have been included is more specific to the content of this book. AI is a relatively new discipline. In many areas of the field there is still not complete agreement on how things should be done. The references to the source literature guarantee that students have access, not just to one approach, but to as many as possible of those whose eventual success will need to be determined by further research, both theoretical and empirical.

Since the ultimate goal of AI is the construction of programs that solve hard problems, no study of AI is complete without some experience writing programs. Most AI programs are written in LISP, PROLOG, or some specialized AI shell. Recently though, AI has spread out into the mainstream computing world, AI programs are being written in a wide variety of programming languages. The algorithms presented in this book are described in sufficient detail to enable students to exploit them in their programs, but they are not expressed in code. This book should probably be supplemented with a good book on whatever language is being used for programming in the course.

This book would not have happened without the help of many people. The content of the manuscript has been greatly improved by the comments of Srinivas Akella, Jim Blevins, Clay Bridges, R. Martin Chavez, Alan Cline, Adam Farquhar, Anwar Ghutoum, Yolanda R. V. Guha, Lacy Hidden, Ajay Jain, Craig Knoblock, John Laird, Clifford Mercer, Michael Newton, Charles Peirce, Robert Rich, Sieve Shafer, Reid Simmons, He Simon, Munindar Tarnaba, David Touretzky, Mazurka Vucoso, David Wroblewski, and Marco

Special thanks to Yolanda Gil and Alan Cline for help above and beyond. Yolanda kept the project going under desperate circumstances, and Alan spent innumerable hours designing the cover and bringing it into the world. We thank them for these things and much, much more.

Linda Mitchell helped us put together many draft editions along the way. Some of those drafts were used in actual courses, where students found innumerable bugs; for us, we would like to thank them as well as their instructors. Thank Mitchell and Jean Scholtz. Thanks also to Speray for his help in producing the cover.

David Shapiro and Andrew Murphy deserve credit for superb editing, and for keeping us on schedule.

We would also like to thank Nicole Vecchi for her wisdom and patience in the world of high resolution printing. Thanks to David Long and Lily Mumrert for pointing us to the right fonts.

Thanks to the following reviewers for their comments: Tegal Areas, University of Southern California; Jai Carbonell, Carnegie Mellon University; Dyer, University of Wisconsin, Madison; George Erns. Case, Western Illinois University; Pat Lingloy, University of California, Irvine.; Brian Schmid, University of Michigan; and James Slagle, University of Minnesota.

Carnegie Mellon University and MCC provided us the environment in which we could write and produce this book. We would like to thank our wives, particularly Jim Barnett and Masan. Mirka, for putting up with us while in the middle of this book instead of doing the other thing, we were supposed to be doing.

*Elaine Rich
gala Knight*

PART I

PROBLEMS AND SEARCH

, WM C ..,...-•-... . . • : . .t 1 : l r . al .t 1 .t 1 .z. 7 **IllaililageMiri:Fl .6. Illreigib-EMI · 11MIMPIMMIMEMWIPMF**

11111111X1110-1Wim

What exactly is artificial intelligence? Although most attempts to define complex and widely used terms precisely are exercises in futility, it is useful to draw at least an approximate boundary around the concept to provide a perspective on the discussion that follows. To do this, we propose the following by no means universally accepted definition. Artificial intelligence (AI) is the study of how to make computers do things which, at the moment, people do better. This definition is, of course, somewhat ephemeral because of its reference to the current state of computer science, and it fails to include some of the very large impact namely problems that cannot now be solved well by either computers or people. But it provides a good outline of what constitutes artificial intelligence, and it avoids the philosophical traps that dominate attempts to delude the meaning of either *first: AI* or *hyperintelligence*. Interestingly, though, it suggests a similarity with philosophy at the same time it is avoiding it. Philosophy has always been the stony of those branches of knowledge that were so slowly understood that they had not yet become separate disciplines in their own right. As fields such as mathematics or physics became more advanced, they broke off from philosophy. Perhaps if AI succeeds it can reduce itself to the empty set. As of now this has not happened. There are signs which seem to suggest that the newer offshoots of AI together with their real world applications are gradually overshadowing it. As AI migrates to the real world we do not seem to be satisfied with a computer playing a chess game. Instead we wish a robot would it oppose to us as an opponent, visualize the real world and make the right moves in this physical world. Finally, the definitions of AI to a greater extent, as we read on, there will be always that lurking feeling that the definitions propounded so far are not adequate. Only what we finally achieve in the future will help us propound an appropriate definition for AI! The feeling of intelligence is a mirage, if you achieve it, it ceases to make you feel so. As somebody has aptly put it – AI is Artificial Intelligence till it is achieved: after which the acronym reduces to *Already Implemented*.

(1469-1527i, Italian diplomat, politica]. philosopher,
musician, poet and playwright

Implemented.

One must also appreciate the fact that comprehending the concept of AI is an aid in understanding how natural intelligence works. Though a complete comprehension of its working may remain a mirage, the very attempt will definitely assist in unfolding mysteries one by one,

1.1 THE AI PROBLEMS

What then are some of the problems contained within AI? Much of the early work in the field focused on formal tasks, such as game playing and theorem proving. Samuel wrote a checkers-playing program that not only played games with opponents but also used its experience at those games to improve its later performance. Chess also received a good deal of attention. The Logic Theorist was an early attempt to prove mathematical theorems. It was able to prove several theorems (from the first chapter of Whitehead and Russell's *Principia Mathematica*). Gelernter's theorem prover explored another area of mathematics; geometry. Game playing and theorem proving have shown that the property that simple machines do them well are considered to be displaying intelligence. Despite this, it appeared initially that computers could perform well at those tasks simply by being fast at exploring a large number of solution paths and then selecting the best one. It was thought that this process required very little knowledge and could therefore be implemented easily. As we will see later, this assumption turned out to be false since no computer is fast enough to overcome the combinatorial explosion generated by most problems.

Another early foray into AI focused on the sort of problem solving that we do every day when we decide how to get to work in the morning, often called commonsense reasoning. This involves reasoning about physical objects and their relationships to each other (e.g., an object can be in only one place at a time), as well as reasoning about actions and their consequences; i.e., if you let go of something it will fall to the floor and break. To investigate this sort of reasoning, Newell, Shaw, and Simon built the General Problem Solver (GPS), which they applied to several commonsense tasks as well as to the problem of performing symbolic manipulations of logical expressions. Again, no attempt was made to create a program with a large amount of knowledge about a particular problem domain. Only simple tasks were selected.

As AI research progressed and techniques for handling larger amounts of world knowledge were developed, some progress was made on the tasks just described, and new tasks could reasonably be attempted. These include perception (vision and speech), natural language understanding, and problem solving in specialized domains such as medical diagnosis and chemical analysis.

Perception of the world around us is crucial to our survival. Animals with much less intelligence than people are capable of more sophisticated visual perception than artificial machines. Perceptual tasks are difficult because they involve analog (rather than digital) signals; the signals are typically very noisy and usually a large number of things (some of which may be partially obscuring others) must be perceived at once. The problems of perception are discussed in greater detail in Chapter 21.

The ability to use language to communicate was another gritty goal. It is perhaps the most important thing that separates humans from the other animals. The problem of understanding spoken language is a perceptual problem and is hard to solve for the reasons just discussed. But suppose we simplify the problem by restricting it to written language. This problem, usually referred to as natural language understanding, is still extremely difficult. In order to understand sentences about a topic, it is necessary to know not only a lot about the language itself (its vocabulary and grammar) but also a good deal about the topic so that unstated assumptions can be recognized. We discuss this problem again later in this chapter and then in more detail in Chapter 15.

In addition to these mundane tasks, many people can also perform one or maybe more specialized tasks in which carefully acquired expertise is necessary. Examples of such tasks include engineering design, scientific discovery, medical diagnosis, and financial planning. Problems that can be solved in these domains also fall under the aegis of artificial intelligence. Figure 1.1 lists some of the tasks that are the target of work in AI.

What is Artificial intelligence?

A person who knows how to perform tasks from several of the categories shown in the figure knows the necessary skills in a standard order. First, perceptual, linguistic, and commonsense skills are learned. Later (and of course for some people, never) expert skills such as engineering, medicine, or finance are acquired. It might seem to make sense then that the earlier skills are easier and thus more amenable to computerized duplication than are the later, more specialized ones. For this reason, much of the initial AI work was concentrated in those early areas. But it turns out that this naive classification is not right. Although expert skills require knowledge that many of us do not have, they often require much *less* knowledge than do the more mundane skills and that knowledge is usually easier to represent and deal with inside programs.

Mundane Tasks

- Perception
 - Vision
- Speech
- Natural language
 - Understanding
 - Generation
 - Translation
- Commonsense reasoning
- Robot control

Formal Tasks

- Games
 - Chess;
 - Backgammon
 - Checkers -Go
- Mathematics
 - Geometry
 - Logic
 - Integral calculus
 - Proving properties of programs

Expert Tasks

- Design
- Fault finding
- Manufacturing planning
- Scientific analysis
- Medical diagnosis
- Financial analysis

Fig. 1.1 Some of the Tasks Domains of Artificial Intelligence

As a result, the problem areas where AI is now flourishing most as a practical discipline (as opposed to a purely research one) are primarily the domains that require only specialized expertise without the involvement of commonsense knowledge. There are now thousands of programs called *expert systems* in day-to-day operation throughout all areas of industry and government. Each of these systems attempts to perform, or potentially all, of a practical, specific problem that previously required human expertise, in Chapter 20 we examine several of these systems and explore techniques for constructing them.

Before embarking on a study of specific AI problems and solution techniques, it is important at least to discuss, if not to answer, the following four questions:

1. What are our underlying assumptions about intelligence?
2. What kinds of techniques will be useful for solving AI problems?
3. At what level of detail, if at all, are we trying, to model human intelligence?
4. How will we know when we have succeeded in building an intelligent program?

The next four sections of this chapter address the questions, following that is a list of some AI books that may be of interest and a summary of the chapter.

1.2 THE UNDERLYING ASSUMPTION

At the heart of research in artificial intelligence lies what Newell and Simon (1976) call the *physical symbol hypothesis*. They define a physical symbol system as follows:

A physical symbol system consists of a set of entities, called symbols, which are physical patterns that can occur as components of another type of entity called an expression (or symbol structure). The symbol structure is composed of a number of instances (or tokens) of symbols, related in some physical way (such as one token being next to another). At any instant of time the system will contain a collection of these symbol structures. Besides these structures, the system also contains a collection of procedures that are used to produce other expressions; these are the operations of the system. A physical symbol system is a machine that processes information through the evolution of the symbols. (Newell, Shaw, & Simon, 1972, p. 110)

Then, the hypothesis is

The Physical Symbol System Hypothesis: A physical symbol system can exhibit intelligent behavior as measured by standards of human performance.

This hypothesis is only a hypothesis. There appears to be no way to prove or disprove it on logical grounds. So it cannot be subjected to empirical validation. We may find that it is false. We may find that the evidence says that it is true. But the only way to determine its truth is by experimentation.

Computers provide the perfect medium for this experimentation since they can be programmed to simulate any physical symbol system. The ability of computers to serve as arbitrary symbol manipulators was noticed very early in the history of computing. Lady Love Lovelace made the following observation about Babbage's proposed Analytical Engine in 1842,

The operating principle of the engine even extends into the domain of any subject to operate upon (although of course no result could then be developed). Again, it might act upon other things besides numbers, were objects found which were mutually in relations could be compared by the abstract science of operations, and which should be also susceptible of adaptations to the operation of the engine. Supposing, for instance, that the fundamental relations of pitched sounds in the science of harmony and of musical composition were susceptible of such expression and adaptations, the engine might compose elaborate and effective pieces of music of any degree of complexity or extent. (Lovelace, 1842, p. 17)

As it has become increasingly easy to build computing machines, so it has become increasingly possible to conduct empirical investigations of the physical symbol system hypothesis. In such an investigation, a particular task that might be regarded as requiring intelligence is selected. A program to perform the task is proposed and then tested. Although it is not always completely successful at creating programs that perform

all the selected tasks, most scientists believe that many of the problems that have been encountered ultimately prove to be surmountable by more sophisticated programs than we have yet produced.

Evidence in support of the physical symbol system hypothesis has come not only from areas such as game playing, where one might most expect to find it, but also from areas such as visual perception, where it is more tempting to suspect the influence of subsymbolic processes. Moreover, subsymbolic models (for example, neural networks) are beginning to challenge symbolic ones at such low-level tasks. Such models are discussed in Chapter 15. Whether certain sub-symbolic models conflict with the physical symbol system hypothesis is a topic still under debate (e.g., Smolensky 1988). And it is important to note that even the success of subsymbolic systems is not necessarily evidence against the hypothesis. It is after all possible to accomplish a task in more than one way.

One interesting attempt to reduce a particularly human ability, the understanding of jokes, to a process of symbol manipulation is provided in the book *Mathematics and Humor* [Naos, 1980]. It is, of course, possible that the hypothesis will turn out to be only partially true. Perhaps physical symbol systems will prove able to model some aspects of human intelligence and not others. Only time and effort will tell.

The importance of the physical symbol system hypothesis is twofold. It is a significant theory of the nature of human intelligence and so is of great interest to psychologists. It also forms the basis of the belief that it is possible to build programs that can perform intelligent tasks now performed by people. Our major concern here is with the latter of these implications. Although, as we will soon see, the two issues are not unrelated,

13 WHAT IS AN AI TECHNIQUE?

Artificial intelligence problems span a very broad spectrum. They appear to have very little in common except that they are hard. Are there any techniques that are appropriate for the solution of a variety of these problems? The answer to this question is yes, alert we. What, then, if anything, can we say about those techniques besides the fact that they manipulate symbols? How could we tell if those techniques might be useful in solving other problems, perhaps not traditionally regarded as AI tasks? The rest of this book is an attempt to answer those questions in detail. But before we begin examining closely the individual techniques, it is enlightening to take a broad look at them to see what properties they ought to possess.

One of the few hard and fast results to come out of the first three decades of AI research is that *intelligence requires knowledge*. To compensate for its inherent powerlessness, intelligence is indispensable. It is knowledge that gives intelligence some of its characteristic properties, including:

1 It is voluminous

- It is hard to characterize accurately.
- It is constantly changing
- It differs from data by being organized in a way that corresponds to the ways it will be used

So where does this leave us in our attempt to define AI techniques? We are inclined to conclude that an AI technique is a method that encodes knowledge that should be represented in such a way that;

41. The knowledge captures generalizations. In other words, it is not necessary to represent separately each individual situation. Instead, situations that share important properties are grouped together. If knowledge does not have this property, enormous amounts of memory and updating will be required. So we usually call something without this property 'data' rather than knowledge.
- It can be understood by people who must provide it. Although for many programs, the bulk of the data can be acquired automatically (for example, by taking readings from a variety of instruments), in many AI domains, most of the knowledge a program has must ultimately be provided by people in terms they understand.

- it can easily be modified to corroborate; and it reflects the aim in (Wittgenstein's) philosophy.
- It is not in a different way in which it is not complete.
- It is not in a different way in which it is not complete.

Although AI techniques cannot be designed in a civil manner with the help of AI problems, there is some logic that is independent of the problems and the techniques. It is possible to solve AI problems without using AI techniques (although we suggest that this is not likely to be a good thing). And it is possible to apply AI techniques to the solution of non-AI problems. In order to try to characterize AI techniques, or to find a way to solve them, we need to consider very different problems, and a series of approaches for solving each of them.

1.3.1 Tic Tac Toe

In this section, we present a series of three problems. Try to play tic-tac-toe. The program in this section is intended to be a simple introduction to the problems.

- Their use of the complexity
 - Their use of the complexity
 - the use of their knowledge
- Thus, they move toward the representation of what we call AI

Program

Data Structures

A nine-by-nine board representing the board. The elements of the vector correspond to the board.

1	2	3
4	5	6
7		

The vector contains the value 0 if the cell is empty, 1 if it is filled with an X, or 2 if it is filled with an O.

A large vector of 81 elements (3⁹), each of which is a nine-element vector chosen specifically to represent the board.

The Algorithm

To make a move, the following is followed:

1. The board is represented by a vector of 81 elements (3⁹). The vector is converted into a decimal number.
2. The vector is converted into a decimal number.
3. The vector is converted into a decimal number.

Conclusion

This program is efficient in terms of time. And, in theory, it could play an optimal game of tic-tac-toe. But it is not.

- li take...s ;ri.)t ot spat in stor12- the tobli.! thin specifies the corro...t move 14.1 make from each board prohibition.
 - **S0111C011C** w illllLi•C to do a 1oE of workmx....cifying ul I thi2 enEricy. in (hie. TnnycLihk.
- .11. i. vcny unlikely that all the required. ninvetable enraicis ari hi &Neirained 'act chilacired withilot ernprN,
- want to eXterlii g,11111,'SLlyLC) lhrO,' dimensions. 'At wouttl Ira\ e (13.43:Lirt fr(¹ and i ai iPaC this techniui: willId rig) longer V hrl tai all. since 3 positions viiot.114.1 io thus 47Vcrwlv]mi li,priosen ntemlyries.
- tecimique embodied ill. (his progn.iitl riteci an) of our requiruttienth for a :12 uod Al
Let `sNee it' wt. Ctirl eifi) flutcr.

Progrant7;)

Data Structures

Hoard A Rim:Li:km:IA vize[or representing die board, a. EIL:seribix] for Provairi L. [4w inNtLtihd. of 1,Aing the nu inheN IL I, or 2 in each elem.!: nt, we 2 findicalin.r blank), 3 iine,,Iii..mtiniz Xt.

Turn tir 5 0).
Anir[LITer move or thE game. i alitaut [o h p]ayed; I irdirates arit: First me,,ne,

The Algorithm

Thumoil) algoriairri [a **SOS** Sithpruccdirreq:

Make 2 Retuirri.Li 5 if lid: center of the hoard is blank. thal is;. Board[5] = othorv.Ise, **Cic** TOLIMk; ;thy blank noncorrwr square (2. 4. 6.. or 87.

Noss %imp) RclMld if player p can ni.11. ',yin (Sri Li !Re). otiverob tufri ihe numl)er of 111e square tlial constitutes nime. Thi% runction eni.ible the proTrani

i17;1rtirl io Hock the c3pronim['n l'LPsswiir 1.1.1>eraleti b citeckirT.. tine zit a i i 1 IH. each of the rows. co.ltimase and diag,onals. Beciluse thlo wxy Ya.14.11:Nare nuilitwred. it {-antcht 311c2ntire row (column or diagonal) to Nee if it iN possihile hy the yalucs. of its,L;111021105 together. Lithe product is IS13 3 2}. !hell X cox win. Jr thk. produci is 50 (5 x 5 x 2). then 0 find ;,1 ii nin rwA, cielerrnine which elerneiro D)1..11111,;. And rimirri 1 fit minitcr of 111,;0 square.

rvtiike% 11 mow. in \L1u:1.mff111ki preiCedilliC^ESea% 11 co: RA kr I ti.) 3 it Turn is tvr 5 IF Torn i even, incromimtii by one.

has. ctraileg3.. far eat :h move it maw havt,;. tc i niake. I I EiLike...[be oe,11d-riumb12-rei.1

the algorithm mores if it is playing X, Ow evr...0rminbered mi.' 'es if it ic [Ailying (..). The; strategy for tan is as fo

I 1 (upper left oorniNl_

1turn-=2 If Board151 ■N (Jo(5),, else Lio(1). Boarci[] i4 Malik.. (01.9), else (r (31,

If 1-30...s.w in I X) is nix 0. then (1c)(Tios2.,winl X J [i _2, h lock oppone.nt's witkb-121,c. Go(rwlake2).

Torn=.5 11Posswin(X) i5 no! L P tl ai ri Go'Posswini X I) ILL., The Poss.witil 00 is not O.. IhIli (it3(Pos:sliviiiL0)) block wit], clam lic)ar(11:11 is blank_ then GLii70., else Gia(3).. [lien:- the prograrn is trying to make a lurk.]

```

Turn=6      If Posswin(0) is not 0 then Go (I)osswin(0)), else if Posswin(X) is not 0, then
            Go(Posswin(X)), else Go(Make2).
Turn=7      Pos;swin(X) is not 0 then Go(Posswin(X)). eke if Posswin(0) is not 0, then
            Cio(Posswin( )). else go anywhere that is blank.
Turn=8      If 1aosswin(0) k nut it then Go(13osswin(0)). eke if Posswin(X) is riot 0, then
            o(PosswiniXA. else go anywhere that is blank,
Turn=g      Same ES Turn =7.
    
```

Comments

This program is not quite as efficient in term, of time a the first one since it has to check several conditions before making each move. But it is a lot more efficient in terms of space, It is also a lot easier to understand the program's sirategy or to change the mategy if desired_ Hu: the toml straegy has still been figured out in advance by thiz programmer_ Any hugs in The pmgrammer's tictarztuc phiying 'Will S11(11W up in the puagrarnss play_ And we Mill cannol generalize any of the program's knowledge to a different domain, such sis three-di IL tic-tac-toe..

Program 2'

This program is idEmii,..4.11 ELL Program 2 excer row one *=hinge in th reprusientation of the baud. We again represent the board a a 111[0-e]enleni vector, but this time we assign board positions to vector demonic. as follows:

3	4
5	9
6	
	2

Notice that this numbering of the board produces a magic square: all the rows, columns., and diagonals sum up to 15_ Thig means that we can simplify the proces or checking fora pcKsible win_ in addition to marking the hoard as moves sire made, we keep a limp for each player, of the squares in which he or she has played. To check fur a pcissi hie. win. ror ono player, we consider each pair of squares. owned by that player and compute the difference between 15 and the sum of the two squares. this difference is not positive or if it is gutter than 9, then original two squares were not collinear and o can be ignored. Otherwise., if the square representing the difference is blank, a move there will produce a win. Since no player can have more than to squares at a :ime, there will be many fewer squares examined using this scheme than there were using the more straightforward approach o f Program 2. This show,;, bow the choice of representation can have a 1119.i0f impact on the efficiency of a problem-solving program,

Comments

This comparison raisec an iniereRiing question about the relationship between the way people solve problems and the way cornpuiers do. Why do people ind !lie raw-scan approach easier while the nurriber-couriting approauh is rflorE..2 efficient for a computer? We do ran know 'enough about how people work to answer that qu.estion completely. One part of the aris.wer is that people air parallel processors and ean took at several parts of the and at once, whereas the conventional complier must look at the squares one at a time. Sometimes an investi alit r o In how people suave. prublerns sheds grcat light on how compuctrs should do so, At other times, the differences in the hardware of the two seem so great that different strategies seem best. As wc learn more alvtit problem solving both by people and by machines+ we may know better whether the same representations and algorithm are best for both people and rrlachinv.s. We wili discuss this question further in Section 1.4.

Program. 3

Data Structures

Board Position A structure containing a nine-element vector representing the board, a list of board positions that could result from the next move, and a number representing an estimate of how likely the board position is to lead to an ultimate win for the player to move..

The Algorithm

decide on the next move, look ahead at the board positions that result from each possible move.. Decide which position is best (as described below), make the move that leads to that position, and assign the rating of that best move to the current position_

To decide which of a set of board positions is best, do the following for each of them:

1. See if it is a win. If so, call it the best by giving it the highest possible rating_
2. Otherwise, consider all the moves the opponent could make next. See which of them is worst for is (by recursively calling this procedure).. Assume the opponent will make that move. Whatever rating that move has, assign it to the node we are considering.
3. The best node is then the one with the highest rating..

This algorithm will look ahead at various sequences of moves in order to find a sequence that leads to a win. It attempts to maximize the likelihood of winning, while assuming that the opponent will try to minimize that likelihood_ This algorithm is called minimax procedure. and it is; described in &Lila in Chapter 12_

Comments

this program will require much more time than either of the others since it must search a tree representing all possible move sequences before making each move.. But it is superior to the other programs in one very big way: It could be extended to handle games more complicated than tic-tac-toe, for which the exhaustive enumeration approach or other programs would completely fail apart.. It is also augmented by a variety of specific knowledge about games and how to play them.. For example, instead of considering all possible next moves, it might consider only a subset of them that are determined, by some simple algorithm, to be reasonable. And, instead of following each series of moves until the player wins, it could search for a limited time and evaluate the merit of each resulting board position using some static function,

Program 3 is an example of the use of an AI technique, For very small problems, it is sufficient. In a variety of more direct methods. However, it can be used in situations where those methods would fail.

13.2 Question Answering

In this section we look at a series of programs that read in English text and then answer questions, also stated in English, about that text. This task differs from the last one in that it is more difficult now to describe formally and precisely what our problem is And what constitutes correct solutions to it For example, suppose that the input text were just the single sentence

Russia invaded Czechoslovakia.

When either of the following question-answering dialogues might occur (and in fact did occur with the POLITS program [Carbonell, 1980]):

Dialogue I

Q: did R.w.,siiiL kin tilik?
A: Russia the amt.il could %kI.: control ilf Czechoslovakia by senclinv truitaps.
V..hat sha uld. the lirlitu-d.
A: The 1....ni1cm.1 Stati:.. should inien.ene

Dialogue ...

Q: Why did Russia do this?
A; BcvaW,L: a wanic.d to incrcce its ipiluenci: over CzechoFil.o:...u.kia.
Q: What should rhe Starm fluir!
A The r1 il tbJ Smes should (leninmnce the. fitis,jan action in the United Nations.

lin lit; POLITICS program., {arts t L:ori:Lititcted by considering both iniput text arid a 2,0parall.12
mociel of the beliefs and actions of various political entities. including Rli]SS:la, Whin th mode] is changed. as
it was between I hese two dial ogu system% answers also ch[i:t * ,lit this example.. the first 4li aloi,Fue
produced whin PMATICS a wade! that was. imorkled to corrspond 1.o !he helicfs or a iypical
Arrierit'ari coriser....alive (circuit 1917). The. • iecon.d dialogue occurred when POILCTICS tivu G 4erk a model that
wa.y inieillkkd to corcspcind to the beTicifs of a 11).02;11 American liberal Rif the same` vii.thigc).

T]ir4Fia L:ri I point her is. that cktning whai i r if) pr odtwo CE)Fterr sv..r.n' ti l-a questiion may ix very
hard_ quintion-answering Li-rota:Ens. donne what it [T]Ctir!-; [t' be an ativiver by the procedure that is
used to 4:olnpu1c [he. ;ins...loci. TI on [heir lilihiOIN In people Ica aLirce tli<<t ml anyivers Enrand by ith:
program 111;11: allo modd 0. defined This

comprolLv saiihVactory.. but no bintur 4.11 cid'wing the problem ha2, yot bon found, For lack of a Niter
we will do the male heck:.. and illustrate three definitions of question answerina; each with
cc)rresponding program drat imple.mcnti
order to he able compare the three programs, we il[tEs[raie all of them using thv roliowing

Miry v. d. 111) pring for a new coat_ S 1112 14 miriLi a (J is r 4rcuLi **Ntic poi it Nth rc. slie ili. rovcrai that**
rfr•l reel ly with her fa'nritedres...

We will also answer each of the following, questions %filth each pro2n.un:

Qh What did Man go !.11-topping ror?

Q2: 'Mar Mar?: find that

Q3: Di4.1 Mary billy anphin2?

Program i

This program interrupts V answer question.; using the literal input text_ It simply moves bus to.txt this men 1_s in the questifins against the input text.

D151 to Structures

ues[jionPutiorns A set of t-rnplsies thus match ccwrimon gucs.rion forms and proilike t tic nt to twzd
match 4.1.gDirist inputs, Templates. and patterns (which vi.ri cal] lest pallerizsi are
that if a template- matches s.k.Jucesst111.14 ail input quem.i]5n then its ii.y.ii:peilitc(1 tee

For example., if the template
 matched against
 the amower to thc. toesricui.
 The input ...timed simply as a itlog cliaracter string..
 stored u a character string..

The Algorithm

To ansiiver a question., du ihe

I _ Compare each el Line 4.1tQuostimirriierro;np..Diml tht~ it t i ;Indtiso thom:thatmall:11
 to genierato a Ott of iext patierrri.

each of paiienis through a substicution that 12..eilerdles. alternative form, of oro
 for example.. "go" in i on vinigh rmini.ti 'Acorn' in the 1r'.x Tilis.step gc,, ne es a new, ex panda!
 set of text patrerns.

Apply each of these text pill:n'1s to Text., and calla:ft all din recalling answers.

4. Reply with the s121 cif unmwew..... just

Examples

Q1: The ierriplato "What slid .1: I rr<a1 c-hes this -and..011.erates the- text paticril "Mari go shopping
 for :." Afigr the poth.:...ni-stillsiitutifiri step, ibis pattern is. ex.paw.led to a soli of patterns inclu.ding
 hirkppinu for Z. and "Mary trmit slopriug. for .7. tatter pullers]. matOe.... [Etc
 11274.U.: the program, using a convention 'hat variables rnatch the longest poy.,sibl c string up to a!...endenee
 !.;tich as a period),, assigns z ihcAtte, 'Ea new coat.' which is as the ansvccr.

Q2: Lillie!.6 [] Ie to %Try larve. allowing for the in of the ohji.21:1 of 'land' bc.i.v.,eeil
 Ifw 11)(0[ry int, plum Nile 1 ikcid. the. i rr ram or the word 'really' in Llic. if. lind I
 he illrti

Q3: the queminn cari. hie answered, then reTonse is 'La rod ont_ "
 Since no answer to this question. unlitained in tho lext, 2TISWer Wilr be found.

Commenes

'Ilk approach is clearly inadquu.te to answe.r the kirids of questions rpeople couldLIn[r reading i
 simple teL^E, t_11 its 4.1.LISIAerthe 1:11:FM qUICSIIIOnS i delicately dependent Ofi the exact forth.

v. Ilia qucsilicirrs arc stakal and on the variations thai .1Aiere in arsign 01 it e terririlates hind the
 pailern suhslitutirms tholi system tu:Lcs, rao, 111: slu:or inudeljulic.y til tbk program 10 peribrun ffic tea
 may m4.1kc you. vifon.Ocr how :such 4311 appro.;Nch coula cv.on b prop-NI:AI. This prin.rraEri is sub.stantiall!... 11.traier
 :away from being useful 1.11.ait th- program wt. looko,d n foi tic-tac-roe. Is this just a str.rivinzw.
 some other tr..4:11nicitio i LIII!good in 4.741il pari.141? wa5, yes. but it i 'orth ntolieming
 that tho 111SpriEmram 1111.171Cly triatellhis palterns, percurning .5011nlih.31[Ons. and
 then answers using-s4raistzloro..ard. combiurnions of canned tL,...xt serricriax fro&riricnLs !maw(' by
 matchE.1., i tilv same approach mat is used in o.n.o. 4yI 'he mos.' famous Tirogaim..; ever written-
 ELIZA, willCh wt' In 6.43. BiliL vi).k 'Dud. thc. rest of this sotiticrice uil prog.int..., it cheiold

becorikb.clear that vi 11 we iiwon c he! tom "iirufkiiLintl.....111.gcnce" dile% not indt34.1c i,LS this
 except by u substantial Ntretchlung

Program 2

This program first (..NArverts the irtpril (ex.' into a structured internal form that .1.tterirpis t li capwre the
 meaning of ihe [scrii.cneelq. It](#) .1.1.sti. cum'cos Liurstiooq intl.) that fon'n_ Ir finds answer.; h TrratchinL2 structured
 forms agains! each other..

Data Structures

English Knowledge

A description of the words, grammar, and appropriate semantic interpretations of a large enough subset of English to account for the input texts that the system will see. This knowledge of English is used both to map input sentences into an internal, meaning-oriented form and to map from such internal forms back into English. The former process is used when English text is being read; the latter is used to generate English answers from the internal form that constitutes the program's knowledge base.

Input form.

The input text in a structured form.

StructuredText

A structured representation of the content of the input text. This structure attempts to capture the essential information contained in the text, independently of the exact way that the knowledge was stated in English. Some things that were not explicit in the English text, such as the referents of pronouns, have been made explicit in this form. Representing knowledge such as this is an important issue in the design of almost all AI programs. Existing programs exemplify a variety of frameworks for doing this, and are three important families of knowledge representation systems: production rules (of the form "if .c then y"), and statements in mathematical logic. We discuss all of these methods later in substantial detail, and we look at key questions that need to be answered in order to choose a method for a particular program. For now though, we just pick one arbitrarily. The one we've chosen is a slot-and-filler structure. For example, the sentence "She found a red one she really liked," might be represented as shown in Fig. 1.2. Actually, this is a simplified description of the contents of the sentence. Notice that it is not very explicit about temporal relationships (for example, events are just marked as *past tense*!) nor have we made any attempt to represent the meaning of the qualifiers, "really," It should, however, illustrate the basic form that representations such as this take. One of the key ideas in this sort of representation is that entities in the representation derive their meaning from their connections to other entities. In the figure, only the entities defined by the sentence are shown. But other entities, corresponding to concepts that the program knew about before it read this sentence, also exist in the representation and can be referred to within these structures. In this example, for instance, we refer to the entities *Man*; *Coat* (the general concept of a coat of which *Thing* is a specific instance), *Liking* (the general concept of liking), and *Finding* (the general concept of finding).

Event 2

initiator:	Finding
resource:	Pam
agent;	Alm.
object:	Thing 1

Thing 1

color:	Red
--------	-----

Event.?

instance:	Liking
resource:	Pam
modifier	Maori
object;	Thing .)

Fig. 1.2 A Structured Representation of a Sentence

SlroclQui.....:-tion A structured representation of the content of the document. The structure is the same as the one used in the input text.

Curiveri n Lii inputText into struciured kx'rii using the [knowlizdg.e. cc.](https://www.knowledgdg.e.cc/))ilaincti EngliNhKnow. 11'Liffr.:
CO nSidCririg :Lievcrai different potential structures. for a variel.% (if ri...asons. includirig the. l'aer that Eng,iish
wordi ctiri be arnhip.i.ous..English sinac[ur.in.; can be pruriourtsIT have sevr.al
FICSSibie 2miectionis. Then, to an....wer do. the follivo.i n12:

]. Chinveri the titles.ii.on sTRICLUTed form. again using the knowleagc 01111¹;:lii1Cd in Use
5.1.)11112 SpeCia.] Marker ill the NittlitUrt iiridien.ii2 The. part of tit ari.F.:ture. that should to ucturrtd agih
answer. This 1111ark.er will o tier!., correspond the OCCUTTL'Ilee 1.11a question word like 'who' Or "'Alai")
in the The exud way in ...vchic this rnarl...i4 gel...-. done depends 4.in the form chosen for
Peptesenting StrticilireciText. If n.qlriiAralfil ler sin...14:11..1N, mtcli {LS, rmirker GAF! be
placed in one ot rn.iJrc slois. If a logical sl'..-tcrn is tr.ii.7(1, li we vex, Markers will appear its ri tblcs i ii Lbw
formulas repro semi
lielatch this structured form against Structuredrext.
Rt.turi tlFre Rilskiver those pans. 4.4 ihe to 'Hui timich the sve.rileni oil

Q1: This question is answered straightforwardly with, J_{ricw}
Q1: This. $\{n_c\}$ also is array successfully "u.rcd eclat".
Q3: This one.. that 4li, c...innot be answered.. since there is no &wed remxirr.u.e to it

This approach, i uhsiantliolly more (know]edgc)-bos.ccl Own tlit ht firs! prognan and St) is more 42111'cuive. It can answnr most quL2stinns w. which replies are ountainctl. in the lext. and it is much le!is hrittle than the Tint program with respect to the *exacr* forn i 4.11 the textc iArlid the questions. As w expect. based on our 'knee with the. rl ati e rr) recognition Rlid tic-tac-toe progrirarris. the price we pay for this increased is time. Aleut seurchinv, variouslint 1 41 lso.; EngliNliKnovy., StructurcirThitt)..

On word cir warning is approprixto hem_ Thu problimi 4.1rpoxlacing ;J knowledge basic for English that is pl.merful enough lief handle a wide ram N. FmElish inputs i every k discussed al gri2aler leng,r.h Cli aptc r 1 5. Iri is now recol2ilifoLi knowledve Eriulish akiile is 1'14.11. ...Ric:Litt:etc ill VIIeral enithiC ii pvoltrion to build ihe kind of slrocuired shown lioru_ Adt!itional world %via(which the text .d.cals i oncn to support lexi4.1 andsymartic d.isuorthiguatiort and Elie corroct tuisigrirnent t1 i aluccedeins Li.) pronouns. art' mg other things. For example. lit the text

Mary walked up to him. "Jim, I asked where he

it is not possible to determine whether the word "she" refers to the roles of Customers and Liles people in mores. To see this, voitirast 01.1rIO:1 iiTiteCedent ¹NhC^{IL} that L•Xi kill the correct 1.11c,2..eclelli k: the first occurrence of "she" in the following.

hlary u:to person, asked her if ...lit' iwedk..t.1 mil% help.

In the sin plr case illustrated our CE jat-buyirt example. i L k pc.)!si hie to derive correct answers. io our iirm. o questions without any knowledge.2c about stores or Cikats., Una till: fact that sortie :such additional inibrniation may be neccs!...ar.... it!. Support qllnilit)31 answering. has ILInzady been illustrated by the lailLlre of this

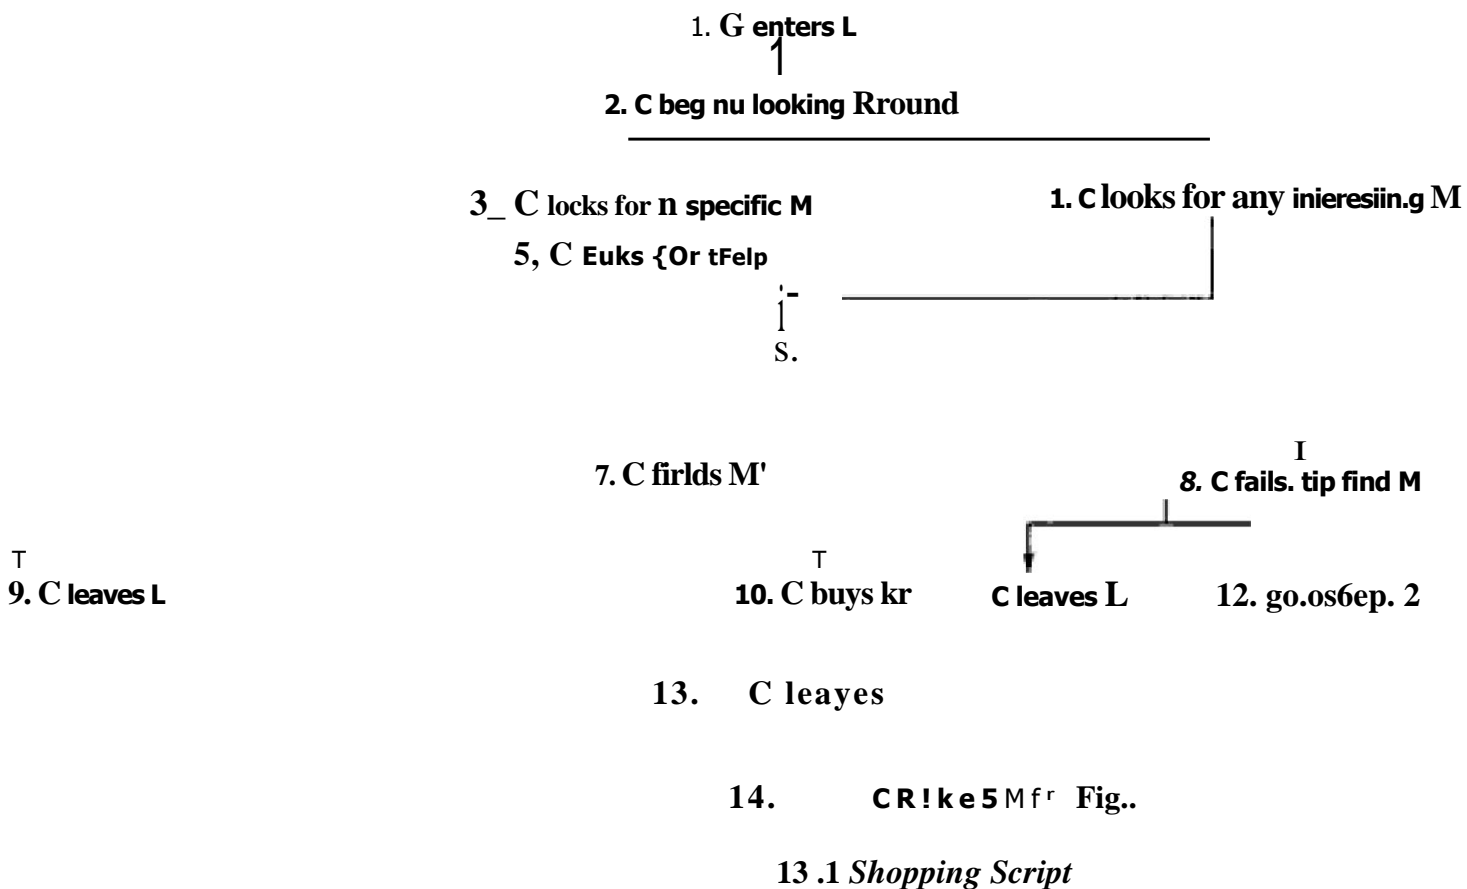
program find. au ...r)mvizr to quizsion 3. Thus we see that although extrading structured repriesemation or Clic meaning of ihr:: input text k au irnprovctnent over the meaning-II= lipprovuzlo of Program 1, it is by nu meangqufficient in general. _ Si:' aced co look at an even mare sophislic-amdl (i.e., know [cd ee-rich) approach: which i4wliai we do next-

Program id

This program converts the input text into a structured form Lh L coilnirri the meanings of ihe wnteng.Les in the tem, and then rt L:Or]lbitleR thaT form with other struo.tured forms [bar des42ribe prior k] at atyout the otijecEs. and situaiirms invc.ilyod in the to I I u.liswors Lluostions using this 'augnitiTted knowledge structure,

Data Structures

Wored.biludcl A structure.(1 representation of hwk.grotind world knowledge. This. structure contain% kiinewle44.e aboui objects. actiOns and situatic.v... r.ha[are &scribed in the input text. This smiciu.n.: is used !{ conn.ruut. Intez,ratulTextit from the input to For example., Figure 1.3 shows an example of a iitstructure hu.1 reivrewt)1...s. symertI¹li knowledge about shopping. Th Li kind of stored k now li dgc shout strcuiotypical exen1L9 is cal L.cd *ij*, *script* and N discussed in mare dutail in St..citi on 10_2_ 'Ube nomiion used here differ from the one norm ally iised in !he li1r2raitire for of 111c prime nolalion describes ;in c.)bjed. of the surrie type as tlic unpr mcd syriiiNil that irnay or may riot refer to the identical object. in the case of oilr text for eximpc, M is a coat and M' is a red coat. Branches in the Figure dcs.criht: attenu.rive paths Ihniugh thL.



13.1 Shopping Script

EieliShicnOW StirT3C is in Fri Tr 2.
Input' rexI The iliput text in charactier tors L

Why couldn't Mary's brother reach her?

with the reply

Bee: dust

But to do so requires knowing that one cannot be at two places at once and then using that fact to conclude that Mary could not have been home because she was shopping instead. Thus, although we avoided the inference problem temporarily by building IntegratedText, which had some obvious inferences built into it, we cannot avoid it forever, it is simply not practical to anticipate all legitimate inferences. In later chapters, we look at ways of providing a general inference mechanism that could be used to support a program such as the last one in the ibis series.

This limitation does not contradict the main point of this example though. Interestingly, this is additional evidence for that point, namely, an effective question-answering procedure must be one based soundly on knowledge and the computational use of that knowledge. The purpose of AI techniques is to support this effective use of knowledge.

With the advent of the Internet and the vast amount of knowledge in the ever increasing websites and associated pages, came the Web based Question Answering Systems. Try for instance the START natural language question answering system (<http://impisart.csail.mit.edu/>). You will find that both the questions – *Whim' is the capital of India?* and *Pr Delhi the capital of India?* yield the same answers, viz. *New Delhi is the capital of India*. On the contrary the question *Are there wolves in Korea?* yields *I don't know if there are wolves in Korea* which looks quite natural

1.3.3 Conclusion

We have just examined two series of programs to solve two very different problems. In each series, the final program exemplifies what we mean by an AI technique. These two programs are slower to execute than the earlier ones in their respective series, but they illustrate three important AI techniques:

- **Search**—Provides a way of solving problems for which no more direct approach is available as well as a framework into which any other techniques that are available can be embedded...
- * **Use of knowledge**—Provides a way of solving complex problems by exploiting the structures of the objects that are involved.
- e **Abstraction**—Provides a way for separating important features and variations from the many unimportant ones that would otherwise overwhelm any process.

For the solution of hard problems, programs that exploit these techniques have several advantages over those that do not. They are much less fragile; they will not be thrown off completely by a small perturbation in their input. People can easily understand what the program's knowledge is. And these techniques can work for large problems where more direct methods break down.

We have still not given a precise definition of an AI technique. It is probably not possible to do so. But we have given some examples of what one is and what one is not. Throughout the rest of this book, we talk in great detail about what one is. The definition should then become a bit clearer, or less necessary.

1.4 THE LEVEL OF THE MODEL

Before we set out to do something, it is a good idea to decide exactly what we are trying to do. So we must ask ourselves, "What is our goal in trying to produce programs that do the intelligent things that people do?" Are we trying to produce programs that do the task the same way people do? Or are we attempting to produce

programs that simply do *the* tasks in whatever way appears easiest? There have been AI projects motivated by each of the goals..

Efforts to build programs that perform tasks the way people do can be divided into two classes.. Programs in the first class attempt to solve problems that do not really fit our definition of an AI task. They are problems that a computer could easily solve, although that easy solution would exploit mechanisms that do not seem to be available to people. A classical example of this class of program is the Elementary Perceiver and Memorizer (EPAM) [Feigenbaum, 1963], which memorized associated pairs of nonsense syllables. Memorizing pairs of nonsense syllables is easy for a computer. It simply replies to them. To retrieve a response syllable given its associated stimulus one, the computer just scans for the stimulus syllable. It responds with the associated syllable. But this task is hard for people, EPAM simulated one way people might perform the task. It built a discrimination net through which it could find logos of the syllables it had seen. It built, with each Wraith's image, a cue that it could later pass through the discrimination net to try to find the correct response image. But it stored as a cue only as much information about the response syllable as was necessary to avoid ambiguity at the time the association was stored. This might be just the first letter, for example.. But, of course, as the discrimination net grew and more syllables were added, an old cue might no longer be sufficient to identify a response syllable uniquely. Thus EPAM, like people, sometimes "forgot" previously learned responses. Many people regard programs in this first class to be uninteresting, and to some extent they are probably right. These programs can, however, be useful tools for psychologists who want to test theories of human performance..

The second class of programs that attempt to model human performance are those that do things that fall more clearly within our definition of AI tasks: they do things that are not trivial for the computer. There are several reasons one might like to model human performance at these sorts of tasks:

1. To test psychological theories of human performance. One example of a program that was written for this reason is PARRY [Colby, 1975], which exploited a model of human paranoid behavior to simulate the conversational behavior of a paranoid person. The model was good enough that when several psychologists were given the opportunity to converse with the program via a terminal, they diagnosed its behavior as paranoid.
2. To enable computers to understand human reasoning. For example, for a computer to be able to read a newspaper story and then answer a question such as "Why did the terrorists kill the hostages?" its program must be able to simulate the reasoning processes of people.
3. To enable people to understand computer reasoning. In many circumstances, people are reluctant to rely on the output of a computer unless they can understand how the machine arrived at its results. If the computer's reasoning process is similar to that of people, then producing an acceptable explanation is much easier.
4. To exploit what knowledge we can glean from people. Since people are the best-known performers of most of the tasks with which we are dealing, it makes sense to use them for clues as to how to proceed,

This last motivation is probably the most pervasive of the four. It motivated several very early systems that attempted to produce intelligent behavior by imitating people at the level of individual neurons. For examples of this, see the early work of McCulloch and Pitts [1943], the work on perceptrons, originally developed by Frank Rosenblatt but best described in *Perceptrons* [Minsky and Papert, 1969] and *Design for a Brain* [Ashby, 1952]. It proved impossible, however, to produce even minimally intelligent behavior with such simple devices. One reason was that there were severe theoretical limitations to the particular neural net architecture that was being used. More recently, several new neural net architectures have been proposed. Those structures are not subject to the same theoretical limitations as were perceptrons. These new architectures are loosely called *connectionist*, and they have been used as a basis for several learning and problem-solving programs. We have said about them in Chapter 1. Also, we must consider that while human brains are

highly parallel devices, most current computing systems are essentially serial engines. A highly successful parallel technique may be computationally intractable on a serial computer. But recently, partly because of the existence of the new family of parallel cognitive models, as well as because of the general promise of parallel computing, there is now substantial interest in the design of massively parallel machines to support AI programs.

Human cognitive theories have also influenced AI to look for higher-level (i.e., far above the neuron level) theories that do not require massive parallelism for their implementation. An early example of this approach can be seen in GPS, which are discussed in more detail in Section 1.5. This same approach can also be seen in much current work in natural language understanding. The failure of straightforward syntactic parsing mechanisms to make much of a dent in the problem of interpreting English sentences has led many people who are interested in natural language understanding by machine to look seriously for inspiration at what we know about how people interpret language. And when people who are trying to build programs to analyze pictures discover that a filter function they have developed is very similar to what we think people use, they take heart that perhaps they are in the right track.

As you can see, this last motivation pervades a great many areas of AI-research. In fact, it, in conjunction with the other motivations we mentioned, tends to make the distinction between the goal of simulating human performance and the goal of building an intelligent program any way we can seem much less different than they at first appeared. In either case, what we really need is a good model of the processes involved in intelligent reasoning. The field of cognitive psychology, in which psychologists, linguists, and computer scientists all work together, has as its goal the discovery of such a model. For a good survey of the variety of approaches contained within the field, see Norman [1981] and Anderson [1985], and Gardner [1985].

1.5 CRITERIA FOR SUCCESS

One of the most important questions to answer in any scientific or engineering research project is 'How will we know if we have succeeded?' Artificial intelligence is no exception. How will we know if we have constructed a machine that is intelligent? That question is at least as hard as the unanswerable question 'What is intelligence?' It can like do anything to make our progress.

In 1950, Alan Turing proposed the following method for determining whether a machine can think. His method has since become known as the *Turing Test*. To conduct this test, we need two people and the machine to be evaluated. One person plays the role of the interrogator, who is in a separate room from the computer and the other person. The interrogator can ask questions of either the person or the computer by typing questions and receiving typed responses. However, the interrogator knows them only as A and B and aims to determine which is the person and which is the machine. The goal of the machine is to fool the interrogator into believing that it is the person. If the machine succeeds at this, then we will conclude that the machine can think. The machine is allowed to do whatever it can to fool the interrogator. So, for example, if asked the question 'How much is 14324 times 73,981?' it could wait several minutes and respond with the wrong answer 1,031.

The more serious issue, though, is the amount of knowledge that a machine would need to pass the Turing test. Turing gives the following example of the sort of dialogue a machine would have to be capable of:

Interrogator:	In the first line of your sonnet which reads "Shall I compare <i>thee</i> to a summer's day" would not "a spring day" do as well or better?
A:	It wouldn't.
Interrogator:	How about "a winter's day." That would scan all right.
A:	Yes, but nobody wants to be compared to a winter's day.
Interrogator:	Would you say Mr. Pickwick reminded you of Christmas?
A:	In a way.

fidinc31a1o1 **Ytt** winieriL4 day. and I do rivil ibink ?Or_ Picloivick would mind
 coiriparistiri.

A: I don't think you'tr serious_ L a 'Ivinter's day or FOLIOS. El typical wintcr's day, rather
 than :1 special one like Chrisin'as_

It will. Ih• j lon.e time beforic a computer pak.!,es. the Turing. SC,Ille p42.1.1p1E' I LL1 is vc none VVer will. BM
 WC arc Willing i• the for less than a complete imitation of a person.. Can wt I1.' eiI klu::thc. achievement
 of AI in Hum! finmEtin.,5?

Often tho Linsi.s'or 14? Lehi s question is. yes_ Sorni2tirries i1 is possible to. get a fairlv precise measure ihe
 a pi41gram. 'For ex ;ill pile. :1 pr4.1griim can ackpiiv Li che in the same as a human
 plo,)]-,r, 1104 rating. B. 1111:rating_5 pho...ers whom din:. program hem_ Alrvigh programs li.,0.12.
 at'Auired L he taiings hi.1-102-L. than aic vast majority of human playon,, jro orti.e.r pr olhlou. dorriairr., a .11275¹.1/4
 precise ineasu•c, of 34 programl`-.a1 h ievement I possible._ Fror ;Ltri p I e. I)EN I)R A L is a program 1 bat
 organic compounds &termiric stn.' LIL Ie_ 1 Cih iIrd III_ 21.'i i1 PaTic irieiLsttre. 4.}i' AL's le% el of
 croirripiied 1 hill-flail chemists, hut it liras pro4litcrEl LinHlyses tlira i lthvo ori,12,inal
 reu11 Thu:, it is certainli2. pierfunning coirily.n.crill).

ill other le 1:1 inic.01 domains. it is possible to compare the time it takes for a program to cti fnpJC tc a tx.ii; to the
 Ii required h yLi per m.l.r) ILI LIB the svirrie thing. For example, there arc v1_11 e rhil [mgrativ, in List by compmer
 courvinicis t rr 4:orifi2Lri.. palOR:War ...o.slems 10 J reeds cot v.11114..11 the pioneer was a progrian called
 R11. 'Thor.: progiirms typically ropin.- ininotes perform laNks that previously required how. 41t a skilled
 c.nr.inecr'N tinic_ Stich priPir.17.1111S art e....alited by 'coking at ithi2 botiorn - wholier tlicy s.ave {or
 makeJ money_

i• marry eiveryday t11011h.itmai,rhe IlarLier Co uneasure a pri..wurri's .rfoirmiaili:e. Suppose,
 fr Wt. u'')(i program ter rrap11nis2 slimy_ For pnthierns such as this, the ho dl. It is
 usually prod,:r'arri respoitcled hi a vdy thai r...rson could llave.

If cur program lirk si1111L1[310 hu]num performance it a tusl.;;_ uL 11)•2 me.al..tire of
 t2:41ioni 10. which tho proFraria's hellaihi i r ci iisrespo rids to that pc rionii [ince_ nwasured h ...a.rious kinds of
 experi •;111dpil1100)1;111i11y .iiILiswe(141nc.11 d Ctr1 }12 vim I that duos kiz, vhtli a i poNsible. We
 • aH 1.1.1112thii L Caik when people do. Nclarious. I. En. el arc] by psychologists. for cornpuriris. individual
 and for ic..stinl,r modol.s can EaLr lased to do this wialysis.

Wc. :Lre forcoil to 1:1,..5ndlock that th4... qui:...stion of 'Li:hotter a machine has intelligelicIZ or can think is too
 netwloto; to .Linswer preci..si:ly, BLit it is often pc);•!•;iblu: to construct a onropulor prugrairi that
 performance. st.i.iirdard. for panicular 'ask_ That ducs ri.ot incaii. dial the pr4i.g..ram don the task in tht best
 possible way. It iineans Pnly that 'owe andersiarid ail least Orlt ti' il:y LA doing at least pan of a task. When ci
 out to design an AI pro,gran], w should attempt to specify as vhc11 possible the cri(eria for success Ilbr that
 particular program lunetioning in its restricted drionaitt For the Criorneri L that k the best 'eke can do,

14 SOME GENERAL REFERENCES

There are great many L,clorce., of information abolii imel]igence. scrme survey The
 broad est arc th.c in alii•vt)lonic Of id... I 9.1] En ; X of
 A rtiffrial Intelligence 11Sheipiro and Edo41th. 0871. hPrith Of which conitlin anicle.s rn each of the major iDpics
 in the field_ Four 43r,,hici IN) 01,,,s 'h al providc gond .c)...cri.ievo, (A the- field are *Artificial lateltigerret. riNinslun.*
198-11, brimiluction 10 Anirr&oi !wen..gowe [Chtriliid]. and rslcDerinoti.. 1985.1. *Logicalf huroi4aridArrs.4.11 Artificial*
[Gene...on2th .inuf Nilsson. ii.M7], urN4.1 Theh'Icnarsh. rthreffignrelTuliimolo, L 1.00C
 more re trio Eel/ scopi2 i c Principe'', Eif_41.1i17.r 4rNik . 9140t which comain_s a CorThal [tcairiberii
 of sonic gonical-put] osc A1 hrii(tu

The higncry resc arch in arri nein! intelligence laNcinati rigsrory, red Liiet.1 by Pamela. McCordi i 11979-1 in her Priachfilre..11,1w Think. Beci1.144: 41111141s1 i.1111 of what we call AI hits he ti' dek eloped over he years.ML:TorL11,1clc.was NI conduct her re.sc;irc.11 for the boa. b I ai....tually inicrviewing almost all of the ppeople ...vhosc vrork, wa, influential in forming the field_

^viost cit iht: conducted m Al hati beer' originally reportcl..1 'puma] articles. conference proceedings. or techilic.11. roporis. Hut some of ihe most intere, iing o,r theLe Flpers have later ^pre:Aced rn special:allections published Comploters and Through! [Feigenbaum lin4J Feldman.. 1963] i vory eiirl3. collection of this sort. Litt!' 4.Inc.:s. 111.C1k1L11.! Sinvim ant! Siklowv [1974 SA:hank arid. Colby 11973], Iiiihrom. ankl 11975!.. WatQroo n all Li I ia 1197 F i rod 112.1. 119791, Webber and Nil ssgjr. 11191. I l. H.iitpt:rn 119861, Shroh.L. 119K4. and several uttlors [bat are mcnrri.i.mial in later chu.ptcn; in connection with spocirc. topits_For nev...erAl parudigrio; t1-1c2 hDak .F.Intehattorto h.. r411(• .1.r0ficire [Tnshincpri. P)9SI a good one.

Thr2 major journal of AI research is culled Nimply inielli;ence. To addition, Cogairfre Sejettrue deccord [0 paper:, deu.]ing the overlappinE. areas of psycholo.gy. arid art i intelligence. Ai Magtazine is a E110143 ephemeral. IcAs technical rru.ig.w.cirbo that is published by the Anwrican Assoiziation for Artificial intelligence (A AAI). f.EPE Erperi. 'PER Tronliarrimrs. on SIIWIrefriS, Man tilnel CyhrTiljfIre , /EFE Tarn enstrairmA pri Neural .verwirky sever;il oitier journal!, pilbli.,11 poiper.... a broad per [urn or A application EL Perb;iih.

Si nc I 969. licit: hx... been a major AI con the. hi tcrnati oint CCillforenc....c: on Artificial I rite I igence arCA Iii hdd every two ivears. prooceding.s u t rh se conferences give a good picture. of the work. that was itakine plaice ;it the time. The other imporiliirit AI conference_ held three oat tvery four VCIEDI sialituu 1980.iss11orp,11100k 1.1k: A A Al ;And its proo.hoclings,14.30. ;ire puhlkhed.

In u.ddii,jun 10 iiiiise 2 eneru.I rererenon, them 12-74.isc!...rr whole irray ni papers and hook; dncribing individoaI Al projbas. Rather than trying to lit them all hero, t1u arc refer-J12d to op; 4pppopriale. Ifiri.K.igliouo the this. book.

1,7 ONE FINAL WORD AND BEYOND

What crodusiow; ean we draw from ibis hurried. introduciion to the major questions of AP 'Mc problems are. varied, in[cresiimg. arid hard_ scl]ve them. w will have ti chi! programs anti p rhap almiter u 'tag of human thoulit. Mlt1 .114.11 Elk; ihe best we inn lo iteria Lr amt e cwiri 1X11 ir wi have soived the problems, wind then vie must try in (14.1 so.

1-1riw actually to go about golviny think: problem!, iN the topic tsar the rest or this hoc k. Virre need mizthods. to hdp LIN Mik Ar s serious dilemma:

J I Ain rrii3!;t contain ;.I Lou of knowledge if it is to handle. anythirp bull trivial toy problornis. EhLt al.. the' 4.im 0 u litok rto od":212. Lirows. it becomes. harder mss the u.pppropriate things when needed,so mc,reknotAilledge fl uNt EL^ LLIdo..1 io H ut now Ebert t even mart: knowledge tc a inanN.re, s() V1111112 idUsl 110.;' Jitiled. and r forth_

Our in AI is 10 construe! working. proE,rrarns that solve the problems we are interested in.. Throughout MOM tar !his. book. Wt ftacus. tan [h• design of repro:;eniatitHi rnechliniNms :Lrid algorithms that Carl USCCI by pl'OrFallik; WE solve the problem., We do not sivrid much time discth.sinE the provrarrrrning proc'iss ptquired E ti turn chew EileGigns into working pfogranh. in theory, it does nE 5E Jrrat ter how thih. process is carriol oul ill what language it k Mont, or can What machine the product is run. In practice, of course. it i5 Oflett much eiNier paxioce as program uF.ing of ceJti rather than another. Specifiru.11y, Al prtrgrnrnI1 apa easiest to build using larkpiap:s that Iu e been designed io uppert syrn4.1tie ruttier th:in prirnarik... numeric computatiorn.

For a variety of reasons, LISP has historically been the most commonly used language for AI programming. We say little explicitly about LISP in this book, although we occasionally rely on it as a notation. There used to be several competing dialects of LISP, but Common Lisp is now accepted as a standard. If you are unfamiliar with LISP consult any of the following sources; LISP1 [Winston and Horn, 1989], *Common Lisp* [Hennessey, 1989], *Common LISP* [Wilensky, 1986], and *Common Lisp: A Gentle Introduction to Symbolic Computation* [Fourey, 1989]. For a complete description of Common Lisp, see *Common Lisp: The Reference* [Sleek, 1990].

Another language that is often used for AI programming is PROLOG, which is described in Chapter 25. And increasingly, as AI makes its way into the conventional programming world, AI systems are being written in general purpose programming languages such as C. One reason for this is that AI programs are ceasing to be standalone systems; instead, they are becoming components of larger systems, which may include conventional programs and databases of various forms. Real code does not form a big part of this book precisely because it is possible to implement the techniques we discuss in any of several languages and it is important not to confuse the ideas with their specific implementations. But you should keep in mind as you re-read the rest of this book that both domain knowledge structures and the problem-solving strategies we discuss must ultimately be coded and integrated into a working program. This process will definitely throw more light into real world problems faced in the implementation of AI techniques. It is for this reason we have introduced Prolog to ensure that you do not end up just reading and believing.

AI is still a young discipline possibly in the 50 years it has been achieved as compared to what was expected. However one must admit a lot more has been learnt about it. We have learnt many things, some of which are presented in this book. But it is still hard to know exactly the perspective from which those things should be viewed. We cannot resist quoting an observation made by Lady Lovelace more than 150 years ago:

In considering any Rev; subject, there is frequently a tendency, first, to *overestimate what we find* to be all ways interesting or remarkable; and, secondly, by a sort of natural reaction, to *underestimate the true nature of the case*, when we do discover that it is *not* as we have imagined it to be. Lovelace, 1843

She was talking about Babbage's Analytical Engine. But she could have been describing artificial intelligence.

While defining AI in terms of symbol processing it would only be right for us to inspect the problem of *Symbol Grounding* [Stevan Harnad, IJRO, The Symbol Grounding Problem, Physics, D42, 335-346] and not forget about it while grasping any new concept discussed in this book. Harnad defines the symbol grounding problem citing the example of the Chinese Room [Searle, 1980]. The basic assumption of symbolic AI is that if a symbol system is able to exhibit behaviors which are indistinguishable from those made by a human being, then it has a mind. Imagine such a system subjected to the Turing test in Chinese. If the system can respond to all Chinese symbol string inputs in just the manner as a native Chinese speaker, then it seems (seems) that the system is able to comprehend the meaning of the Chinese symbols just the way we all comprehend our native languages. Searle argues that this cannot be and poses the question — If he (who knows none of Chinese) is given the same strings, and does *exactly* what the computer did (maybe execute the program intuitively!), would he be understanding Chinese? The rhetoric finally leads to one tenuous inference — *The computer does understand Chinese*. It is thus important to note that the symbols by themselves do not have any intrinsic meaning (like the symbols in a book). They derive their meanings only when we read and the brain comprehends it. It goes to say that if the meaning of the symbols used in a symbol system are extrinsic, unlike the meanings in our heads, then the model itself has no meaning. As the symbols themselves have no meaning and depend on other symbols whose meanings are also intrinsic, the reasoning around meaningless entities which is the symbol grounding problem.

In the context of the meaninglessness of the use of symbols, Harnad provides a classic example of learning Chinese. You may not know Chinese and had to learn it using a Chinese dictionary. You

would compare character by character of a given word and find the corresponding word in the dictionary only to find many more (meanings) written in the same language alongside, for which you would repeat the same task. The process would put you on an endless merry-go-round. It would be only by translating it to a language that you understand that your brain can finally perceive what it signifies. The Chinese symbols in the present case are not grounded to its meaning. The moral of the example is simple — *You cannot round the meaning of a symbol with other meaningful symbols.* Hamad 4.11 So does that cryptologists are able to comprehend ancient languages and symbols because their efforts are grounded in their real world domain knowledge as also on some previous language that forms its basis.

Robots form the ultimate test-bed for AI. While AI researchers have brought forth a reasonably large repository of techniques and programs that are based on the symbol system, implementing them on robots have posed several problems. Though this may be beyond the scope of this book we must exercise caution in its implementation in symbolic AI. For instance on board a robot a symbol 'red' has to be actually grounded to its value, e.g. represented by the camera or a colour sensor.

Finally one should not forget that research in AI is multidisciplinary. People have been using AI techniques to reap benefits in a gamut of applications. There are still a lot more untrodden paths to be discovered. In the quest to find better techniques, the reader is advised to give imagination a free run so that the marginal and the peripheral are avoided without losing the grounding of each symbol.

EXERCISES

— IV — 15 — — — — — ANNEXIVE —

1. Pick a specific topic within the scope of AI and use the techniques described in this chapter to do a preliminary literature search to determine what the current state of understanding of that topic is. If you cannot think of a more novel topic, try one of the following: expert system for some specific domain (e.g., cancer therapy, computer design, financial planning), recognizing motion in images, using natural (i.e., human-like) method, for proving mathematical theorems, resolving pronominal references in natural language texts, representing sequences of events, or designing a memory organization scheme for knowledge in a computer system has on our knowledge of human memory organization.
2. Explore the spectrum from simple to AI-based techniques for a problem other than the two discussed in this chapter. Think of your own problem or use one of the following:
 - Translating an English sentence into Japanese
 - Teaching a child to subtract integer;
 - Discovering patterns in empirical data taken from scientific experiments, and suggesting further experiments to find more patterns
3. Imagine that you have been to an aquarium and seen a shark, and an octopus. Describe these to a child who has never seen one. What resources and modules does the child use to comprehend the nature of these marine animals?

CHAPTER 2

PROBLEMS, PROBLEM SPACES, AND SEARCH

!-C fitm' (/rij I:V 2;0 Mart..

*ft dia /slat_y'

FroMwitt kweier

—Albert Einstein

(1879–1955).. German-(t)4.mi thucProlcitl pad Nisi

In this chapter, we gave a brief introduction to the kinds of problems with which AI is typically concerned. This is a collection of examples of the types of problems it offers to solve. To build a system to solve a particular problem, we need to do three things:

1. Define the problem precisely. This includes precise specifications of the initial situation and what a solution is. Analyze the problem. A few very important requirements can have an immense impact on the appropriateness of various possible techniques for solving the problem.

2. Find relevant background knowledge that is necessary to solve the problem.

3. Choose the best problem-solving technique, if any, and apply it (then adapt it to the particular problem).

In this chapter and the next, we discuss the first two and the third of these issues. Then, in the chapter, in Part II, we focus on the issue of knowledge representation.

2.1 DEFINING THE PROBLEM AS A STATE SPACE SEARCH

Suppose we start with the problem statement "Play chess." Although there are a lot of people to whom we could sell this and reasonably expect that they will do as we intended, as our request stands it is a very incomplete statement of the problem we want solved. To build a program that could "Play chess," we would first have to specify the starting position of the chess board, the rules that define the legal moves, and the various positions that represent a win for one side or the other. In addition, we must make explicit the goal of not only playing a legal game of chess but also winning the game, if possible.

For the problem "Play chess," it is fairly easy to provide a formal and complete problem description. It may be described as an 8 × 8 array where each position contains a piece standing for the appropriate piece in the official chess opening position. The goal is to reach a position in which the opponent has no legal move left; that is, the king is under attack. The legal moves provide the way to get from the initial state to a goal state. They can be described easily as a set of rules consisting of two parts: left side, that serve as a filter to be matched against the current board position; and right side that

The eNtre.Trke of this approach Is shown in the first tic•Inc. *toe* prell:rain of Chapter I_ Each (...ritry in
ector Q.oliTosposiici. IL) a rule

art rati4)Jl.1111..Irtside.tifi each 02.Lcri gurat I E1
ctpresented implicitly by th.c. iridcx position_ The right side. of each rule. diz!,cribes. the upc...T.idin.n. to be
perforthod and is reprrsented by nin::elenicini vector that ctErre.spomis ti the
Each of them...! rules is maximally Npecitic, it applies only to .i sinyle Niard. and, vis is result. no

required when Kuch rules are used. However, the drawback to thiK exirerne approach r !hat the
problem solver earl take: ric 4.1.4.:Eitan. at all. in a novel siination. Irr tact, c2.gseritiftlv no problrn 5.49/1.ing really
occurs.. For illiC-C4.14.:toe playing program, (hi.,l I R F d prLiblim, since it i ible criirnerate all the
Nickwion4 tHvrd ions) Oral may 4 11.3.k.I ki. TIM fir Inriom prt11'51121111b;.. this. k nilt Ihq 4.71...e. In orddr to
so] vc nevi. problems_ more .E.cricr:11 inusi

1st second is ex.c..mplified tly rules '3, slid .1 in kr.. 2..3. Shtltilli [1105 Or should they not ki included i r5 tho
)ist or avvilahc operators') Emptying an tutnicastiroi amount ot'v..aio onto the gintind is cc rtainly allowo3 hy
;tic pnibluni sti.i[ernent_ Bul a superficial prolirriilarir analysis. of [13e proble3II 311akE'S it clear doinp eta v ill
clever „Lbt i1111Y dEPNer [0 a solutinn. Again.. we sec. rlc
j1.3S.1 the prObieni ItseIC. 4124 Oppi.X41,'d Li) nil problem ;Ind sonic. kriovilcidge
solution.

Ilu1 11 and 12 &Lisa'aft. a Third issue. To sec to dratthi250 Lepre.seni..
lookat the 1w :it tiao Noluition...ihown in Fig_ (Onee the! ctale (4, 2) is re lied, colivintis h
k'dinne'1111. 121111; produced, but thlev ;tic. in Ow.iik...reconv. St} the thiriL.' IL) (Ii) i
111,rr 3 RIM bcfon: IhNE cHri be 4101112, the water 111;it i Areaci.. Ljm24.2.;illon jug ri51ININ: cilphicd
out (rtlk r). The iLica bellind these speciat-parpose vuks is to capium the 1:now[edge that can kv
usiv,1 this stage the probl2rn_ Thcs.;.. rules do not nctually acki privicr to the s'sti.:rn since the
up21;itic.)rts they describe. are already rim% [clod b) ru.L- 9 the ca,e of rule 1 I and 1'5) rule 5 fin the c.a.s.e of ruk
12). fiut.. dept-nti* on lie ciairol strare:4) that is. ustLi for selerliryi i l to tr y [luring pmbem
OleLhe 111 in.ay klil-gra4,1,... perform:mix. 111.31 irk: use (11 th..2...e. rulc.. riny also improve puTiorrilance ir
preicri.mco is siktialcasc. rules (as viirc discuss. in Seciioi)

We have row discussed two quite different problemn:.. chess and the tivater 11.1.11' problem'. From
cliscus siorlc. it should be dear that the 1)P:it LiterP toward the deslgnora progr:im to solve a problem' must he the
crealiin fit' a formal and manipulable. Lte!..cription of the. prol,11.1m 11.; [Ornately, we 'would like to he uhie to
write programs I E•i a[can ploducc such forinal descripaorp. 111.11111 infL51 Ogle S. Th i ,process is
called f.tpuraffoolrfirNtriorr. h. is no! ot tL 1 lwon -under:;i lww lo SLICT1 progrorm. hmt sce Section
17_1fura(14;cripti or tne prugrui. Lbw pleci: prfIBICT11. Until ft h.:curries possible ID) autuniate.
tills process, it must hi cifinc hi... baud, however_ For :cimplo prohl SLich a che:is he wa[er jug., this is not
very difficult. Tile problem!, art: artificial

h
II' i
]..5111. I) structured. For oitcr problins_ varlicularly
occurrine, tines, [his step is much MON thie tw..1., of vecifying precisely

1121P/q onlersuirid an Eliniglish Szlitence_ such a specification must sonielow provided bfore
we can deLiig.ri a program to solve the problem, producing, such lx.Fcific'ation i itself ven.. hard problem.
.A.h.hough our ultimate goal is to be able to solve difficult. unstrilemred prubleins. such naruiral language
Lifiderstimping, it is t12.1111 IE3 cimpTor problem:_ ...1.14.:11 as the waleriti..4. prolliern, in order to gain, insigto
into tile details of mileillE)LIN tha It u i I I Ionia the bor,k for !hie hardL.I

Summuzing%Oral h;ivc ji.ht said. in order to provide ;1 form..11 (tc.scripitrn of a problem, we rims! do tile
following':

1. Detine stale %pace that contain!, all thy: possillk configurations of ti k,. (511;jc2Ch
44)110 nne.e.), ti'si2, i L ihlt 141 c.kfille this KINI.Ce With() Lrt expl ii:itly eniurniffa [ins all
or UR: SU.NicS ii cuntains.

- 2_ Specify one or more states within that space that describe possible situations from which the problem-solving process may start. These states are called the *initial states*.
- 3, Specify one or more states that would be acceptable as solutions to the problem. These states are called *goal states*.
4. Specify a set of rules that describe the actions. (operators) available. Doing this will require giving thought to the following issues:
 - What unstated assumptions are present in the informal problem description?
 - How general should the rules be?
 - How much of the work required to solve the problem can be done by the rules?

The problem can then be solved by using the rules, in combination with an appropriate control strategy, to move through the problem space until a path from an initial state to a goal state is found. Thus the process of searching is fundamental to the problem-solving process.... The fact that search provides the basis for the process of problem-solving does not, however, mean that other, more direct approaches cannot also be exploited. Whenever possible, they can be included as steps in the search by encoding them into the rules. For example, in the water jug problem, we use the standard arithmetic operations as single steps in the rules. We do not use search to find a number with the property that it is equal to — (4 — 1). Of course, for complex problems, more sophisticated computation will be needed. Search is a general mechanism that can be used when no more direct method is known. At the same time, it provides the framework into which more direct methods for solving subparts or a problem can be embedded.

2.2 PRODUCTION SYSTEMS

Since search forms the core of many intelligent processes, it is useful to structure AI programs in a way that facilitates describing and performing the search process. Production systems provide such structure. A definition of a production system is given below. Do not be confused by other uses of the word production, such as to describe what is done in factories. A production system consists of,

- A set of rules, consisting of a left side (antecedent) that defines the applicability of the rule and a right side that describes the action performed if the rule is applied.¹
- One or more knowledge/databases that contain whatever information is appropriate for the particular task. Some parts of the database may be permanent, while other parts of it may pertain only to the solution of the current problem. The information in these databases may be combined in any appropriate way.
- A control strategy that specifies the order in which the rules will be compared to the database and a way of resolving the conflicts that arise when several rules match at once.
- A rule applier,

So far, our definition of a production system has been very general. It encompasses a great many systems, including our description of both a chess player and a water jug problem solver. It also encompasses a family of general production system interpreters, including:

- Basic production system languages, such as OPS5 [Brownstein *et al.* 1985] and ACT* [Anderson 1983].
- More complex, often hybrid, systems called *expert shells* which provide complete (relatively spanning) environments for the construction of knowledge-based expert systems.
- General problem-solving architectures like SOAR [Laird *et al.* 1987], a system based on a specific set of cognitively motivated hypotheses about the nature of problem-solving.

This current chapter is for the use of and right sides of forward rules. As we will see later, many systems reverse the sides.

1111

Ali of these systems. provide the overall architecture: of a rpnxi action system and allow the prugroi'rnr v..) wri(c rules that iiefine particular problem5 to be solved. We di 5f,,uss priAlunion systemri i...isLIC5 further in Chiüpter ti.

We havi: now ...een chat in order to solve a problem. we must Erse reduce it to one for which a precise SIOIMIII L,;an he gi'err. Tlii-L can be done by delinlng the problem's sratib.space (includinis. ihe mart and Qoul st9 to si anti 21, YL't UT opt:Jailors for moving in t hat 5pace. The problem can then be solved hy sears ii n LI P4 ar ;I path through the spnce from i:kri initial state in. a goal s.state. The pcocesN of solving the problem can userully ht modeled LS a prod LiCtiOrl Systerit 111 the. rest u r th is suction. wo look at the problem of choosing the appropriate control structure for [hi. production soystem _so tha4 IFic search can he Sys. efficient Lis possibt:.

2.2.1 Control Strategies

So far. we have completely igriorcJ the toe:slim of huw to dceide which rule: to apply next. during the process of :searching for a soluijori to a problem, This question arises since often more than one rule (and sometimes rewor than {ng; rule) will fluve its left side Tria;cl'E the current sii.le_ 'Even viüithout a great deal of thought, it is clear that how such ddeisiort, ar12. made will hair'ff a crtwtial inipaci on hinv quickly, and even v...hetber. .I problem is finally solved.

○ The'ji..sir 0elm: woven; of a goad roluoi ..cfrfirregy is !hal it undses moriorr. Consider tigan the water jug problem of the last section, Suppose ...ve implemented the simple L:ontrol straiegy of startling each time al the lop, Or thc. list (II' mks and choosing the lirs.i jr priicahle ont.. If we did tlini. ^...ve wouci never !..wive the problem.. We would continue in filling the 4-gallon j u with vhalwr Control straigr.ies that do. TICK Cii.LISC motion will never lead in a s,nlution_

• The .vecorid requiremem rif a good corrirof ararn;ry is lime il he p....i.reararic. 1-kre is another simple Control smii egy for the water jug prohlerlrir. On each Cycle. chnow at rariclom froin .airicing the zipplcalaie nI In. This 'ti.E Ed c^{Syr} is buyer it' an the Firs. 1 r 4: ancs. nio Lion_ h will lead to a. so] ration eveninall. But we are likely Eta 2irrivu at the szirric slab!: several riprics during EFIC pruce;ss and to 'use man) rntirc sims !ban are necessary. Illocause the control strategy is non. _;ystematic., W1 may explare a particular useless sequence of operators several tirt es. before we finally lincl a .solution. The requirement that a control qr.ategy he systematic¹²⁰ responds (1.1 the need for w:litllr,11 rnmion (over am cr yurs.k... or scir•ierai steps) as well kig fur]sac al rnotiLill 11.11VE:r the curso 0.1' 41, single step'. Cum systmialic cl mtrul ratan :6.y for.hc.. wailer jug proh le ni is the following_ Construct si tree with the initial mate as its rg.mat. Generate 3[[the: offspring of the. root by applying Each of the applicable rules to the initial state_ Fig_ 2_5 shows how itlh.e tress looks at Ehii point. 'Now for each leaf node, gmerate all its successors. by applying all ttic rules that are appropriate. The tree at this pin is shown in Fig. 2.6.⁴ Continue this prose so; until some role prodn:eq [1. goal swill.. This process, caitni *brigirdi.or-frrsir search.*, can be describe.2d precisoly as follows.




Fig. 2,5 One LvveII ola arcoatti-
First -Search Tme




Fig. 2.6 Boat Levey vrr./ Braga -
Firs Serrre...lr Tree

" Rule . 4, I]. 5i12 have ti.Lerr 4;11ored ill cons.tructilip. theLbk.i.puli litu.

Chess Traveling Salesman Tic-Tac-Toe	in a crucial advantage of our side the uponcoming the Number of the distance., Far I for r''win White we could win arid in white [move on: piece plus 2 for e.L.1.4:1] 21.11:13 row in which we have two pieces
--	---

Fig. 2. *Simple Heuristic Functions*

the purpose of a heuristic function is to guide the search process in a most profitable direction by suggesting which path to follow first when more than one is available. The more accurately the heuristic function estimates the true merits of each node in the search tree (or graph, i.e., the more direct the solution process) in the extreme, the better the function would be so that essentially no search would be required. The system would move directly to a solution. But for many problems, the cost of computing the value of such a function would outweigh the effort saved in determining the best move. For example, it would be possible to compute the value of a node in question and determine whether it leads to a solution. In general, there is a trade-off between the cost of evaluating a heuristic function and the savings in search time that the function provides.

In the previous section, the solutions to AI problems were described using a search process. From the dimensionality in this problem should be clear that it could be precisely described as a process that

heuristic search. Some heuristics are used to define the control structure that guides the application of rules in the search process. Others, as we shall see, will be exploited in the rule themselves. In both cases, they will represent either general or specific world knowledge that makes the solution of hard problems feasible. This leads to a way that could define artificial intelligence: the study of techniques for solving exponentially or polynomial time problems, exploiting the knowledge of the problem.

2.3 PROBLEM CHARACTERISTICS

Heuristic search is a very general method applicable to a large class of problems. It encompasses a variety of specific techniques, each of which is part of a set of techniques for a small class of problems. In order to choose the most appropriate method or combination of techniques for a particular problem, it is necessary to analyze the problem along several key dimensions:

- a. Is the problem decomposable into a set of (nearly) independent smaller or easier subproblems?
- Can some steps be ignored or at least undone? b. If they prove unwise?
- Is the problem's universe predictable?
- * Is a given solution the best possible? (Can all other possible solutions be found?)
- * Is the desired solution close to a "goal"?
- * Is a given solution absolutely required to solve the problem, or is knowledge important only to constrain the search?
- * Can a computer that is simply given the problem return the solution, or will the solution of the problem require interaction between the computer and a person?

For each of these questions, I can examine each of those conditions in greater detail. Notice that some of these questions involve not just the problem itself but also characteristics of the situation that is desired and the circumstances under which the solution must take place.

2.3,1 Is the Problem Decomposable?

suppose mu wawa ;0 miye the prob112111 from puri az the ex prc. Si OD

$$P_{-r^1+} + 5irt^{2*}$$

We can solve this problem by down intro ihree smaller problemR, .e.o.ch of which we can 'her FiLilve hy ming it 74114111 4,:41.11ection 01 spet:ific INiFure 2.9 shows the problem tree that will bi gcrien3ied b the-proce SN prOblernnnn_' it Canby exploited

S reCUsiVe in Legration program that 'Works as follow; Au cacti step. it checks to scc whether the priEb-lent it un is inirricdiakly &olvthile. If so, thvin

answer ig returni211 If the problern is not elpiily solvable, the iniqr:irrir checks to Scc' *whether ill earl decompose the problem smaller problems, iii can, it erciinos those. problem ariti coils i i..cel i recurs ive I y n ti'. rn_ Using. this. tvchnique of pre.thioti *decomposi-tion_* we can often stall e. very large problems. eaqlly_

Nov.. coirri(lor the pr4.11.11ow illtisirdtud in Fig. 2.10. This problem is; drawn from the dontu.in often roferrod to in Al. literaurc the *Nadu world_* Assume that follow[ng '1iFeratnirs are available:

1_ CLEAR (Al [1flock has nuihing up and put. it tin au: lahlej

2_ CLEAR ct) Lind CLEAR (y) UN yi rptit x on yl

Applying. lilt lechnique of problem decumpoL,ici)31 to this RimpLe Mocks wo-rli..1. -example lead co' a trc e sucti that shown in Fig. 211 . In rch, goat4 are underlined. Slid e that have aro not underlined. The idea. of this solution is ti.) ccducc the. problem of Exiting B on C and A (311 13 to VI&O :;cp4. a' ate problems. The first of the ncw hlkin getting 13 on C. i simple.. given the start slime.. Simply put B on C. The ...e_conr..1 LIFT02111 is not cluii p I _ Since the only c3pelitl.4.)ts we have us to pick Lip .single blocks at a 'lime. we have lit clear off A hy C before we can pick up A and put ir un B. '17-ii: can easily he duce_ However, if we nov, try io combine the two subsoludons into one solution. We will Regardless of nne we do first, live will no[he .ii.13]12. in do (he second as .1.102 kid planned. this problem., the two 611bproblems am riot independent. They interact and the qe interactions muse be oongidurni in order to arrive at a..6olution. tor the entire problem.

These two exarnplos, symbolic integration arid the. Hocks world, illustrate the. difference betwocn dcompoisuble and nuridocompo.sabk prublems, In Chapter 3, wc pre-sc.:A a specific algorithm for problem decomposition, and in Clicapter 13. we look al. what happens he decomposition is impossible.

2,3,7 Can. Solution Steps Be Ignored or Undone?

Sluivosc iiro trying, ki pVOVeza nialiherrililtical theoPuM. '1/4'1.74! prXcal by p-CIA inga kmnlil that wt. thirik will bc. use Eventually, we realize that (hi:. lernrna is no Ih.clp. at all_ Arc tve in truuble?

$$\begin{aligned} & \int x^2 dx = \frac{x^3}{3} \\ & \int 3x dx = \frac{3x^2}{2} \\ & \int \cos^2 x dx = \frac{x}{2} + \frac{\sin 2x}{4} \end{aligned}$$

Fig. 2.9

.4 Decompagabie Probirem





Fig. 2.11 A Prot] .see Solution far d Mocks Probitr;

2.3.3 Is the Universe Predictable?

suppoN2 that we learn pi ay jug with the S- puz It Every time we :mike. a 11101:1'. 4r know curtly what win happen_ This mean s that it is possible to plan an entire sequence of ma'es and be confident that we know v4 hat the rcsultin.2 state v:1 be, 'Vide can 1.1'Lid: planning to avoid having to undo ;Lau:ill moves. although it will still h neecess.ary th, hac:k'troct past those one at a time during the planniuE. process. Thus a control !Arnow rc that ilkporls 11acktrucki will be

▪ L31 gik1111JS 4111.12V1/1Mfil th i 8-1)141:111.12- h i pli4nning proc-vA' May not 1.112 possible. Sup.p4I've VIII: ti plas' bridge. One of The decisions we will ic) TN:Lk:i which i2 arid to play on the first trick. What we would to dui Is to plan the 'entire. hand before [hulking that first play. I Lit Bow it is foil possibk to do midi planning with cerhaility since We' cimnot loriLm ex:Act!)'he,21f Al the c.iirds arc. what ihe phi. ers; will do on itwir 'urns, The hiest we Call 'Lit) iptVW[gatC vcra pimp.. and use probtahitilits of ou[coriles Liu chcw..c- tt plan 'hat has the highcsi estimatod probability Of lcadng tc a gui.a.d1 scurt!n I b the- hard

Thine illustrate the difference be1...y.1_42-n certain-o titcome le} and 11 riccriui ri-ouworne brjdge problizrns.. Orie way or plarming is thiLit it is problem-solving without rc-c..dhack from the environment_ For solving CC nairl-outeonio problems_ this open-loop approzich %%AI work rim! rcshil[of an acli c.iin be piedktd planni12 LP-12L.Imgchlkhrilk ci]xbrithn, that i trivall teed to lcn d. solution. For 10[14:421rin-orit..21 H1re prc 511)1C UN, liowc.ver,*lining. can it —helm 42..2-11 crate d soqi.pcncet: of opersiors that has a good probability of 1.12...itliiirg to a solution, 'a) oIve Sill:11 problems, vire ilecd for ;1 process of *plan rerfgon to take ploc* the plan k carried out arid the nocnsary feedback iN provided_ In addition to pri vi di ng gtirantev of an ;iota milution. planning iur u.ncerhain-raii [come problemn has1.112 ithed it is (iil'en very expenqive shrhee the number 41.1 soklion pi.olik That need tif h explfired

t MIdJy'OrMit the, nuimtihe` 111f poirns ;it carinot

The labt two problem characteristic!. we have di-ig.i.1;e4.11, iginon.thIc verstli. roc over:Able in'ccoverabil,!. and certain-outconic vcrso!. unceriain-outconle.. into raLi in an in tc.n::....ting. way_ As Dircu.dy been mcntioncd. one way to irreerovurable b terns is. to plan an entire soil ution b• [urn ie.tub:irking on an imp] Li nee tai i on f..)f pkin. But alb. planrling process can only he iluite kiTectively for 6:E.:mini-Lim too no Thus (me o the hurc.le:L.t types Of problem., til solve k M1 e irrr..b.cove.rablth. tinvvrtiiii-ntilcorne. A Fink: cvainplos 01 such proble ms aro;

- Playing. brid2e. hut can kill vhe h; the lable acciaraw csiniutc.... of the probabilities of L:15di of 1 1:14 ixlssible
- COrarolling a robot The tuttcome Iv; uncertain Ibr a 'variety or masoiri.... 54.1i7ic'one might nurs.c:

sometling. in tiro the path or the arm. The :Lcars of th• ann might Mick. crnir LIM Id cause thc. air It to knock over a whole stick of thirigN.

- * lliwyer decide hiya t4.1 deren4 his client against a nutnier ocharge. Herr4 we probably canclut even list all the po...:sible outcomes. much PrOhilbili6CS.

2.3.4 Is a Good Solution Absolute or Relative?

Considur the problem) of iuti...vcriET

Lased on it

such (Fic

. :...1.1.reus was

\kirur:~c7.1s 1³onipeiart.
1%o.147-reu. bry,rjibomin40A.D.

4. Ail mr.fn

5_ All Rompeia_ns died %%hen ii.okano erkirpted in 79

6. 'ii ink hk.in 1.50

7- Ii is now 1'19] A-D.

Suppose we ask the question *Is Marcus a philosopher?* By reasoning in a formal language as predicate and then using formal inference methods we can fairly easily derive an answer to the question.⁷ In fact, either of two reasoning paths will lead to the answer shown in Fig. 113. Since all we care about is the only way to the goal, it is not matter which path we follow. If we do follow one path successfully in the answer, then: is necessary to remark that. *see* if some other path might also lead to a solution.

	von
Man:us was a philosopher.	axiom
4. All men are mortal.	it)M 4
8. Marcus is mortal,	I 4
3. Marcus was born in 40 A.D.	axiom 3
7. It is 1991 A.D.	it is 1991 7
9. Marcus' age is 1951 years.	3. 7
6. No mortal lives longer than 150 years,	
10. Marcus is dead.	S. 15. 9
(JR)	
7. It is 1991 A.D.	it is 1991
5. Philosophers died in 79 A.D.	axiom 5
11. All philosophers are dead now.	7.
2. Philosophers were a philosopher.	axiom 11 2
12. Marcus is dead.	11. 2

Fig. 2.13 Tom ways of *Duricith, I Mot ittorco. Dead*

But now consider a traveling salesman problem. Our goal is to find the shortest route that visits each city exactly once. Suppose the cities to be visited are [from distances, here] shown in Fig. 2.14.

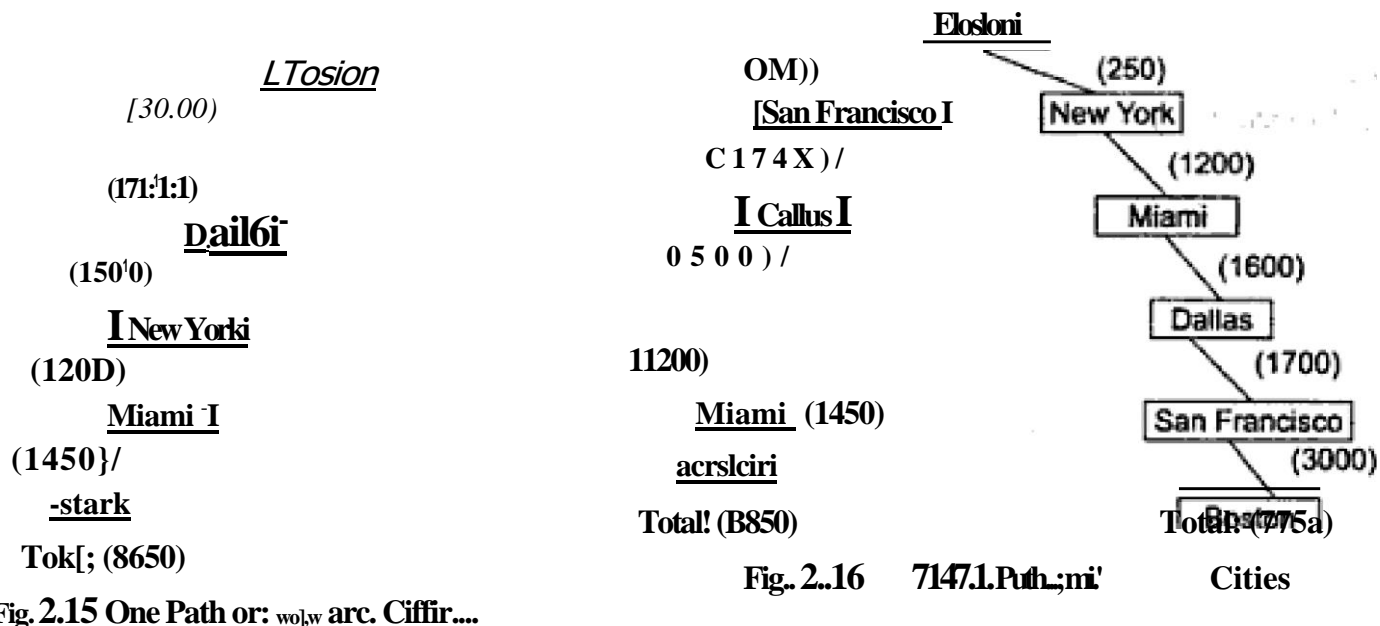
	Boston	New York	Miami	Dallas	S.F.
Boston		250	1450	1700	3000
New York	250		1150	1100	1100
Miami	1450	1150		1100	1100
Dallas	1700	1100	1100		1100
S.F.	3000	1100	1100	1100	

Fig. 2.14 an instance of the Traveling Salesman Problem

Where place the salesman could start is Boston. In that case, one path that might be followed is the one shown in Fig. 2.15 which is 8850 miles long. But is this the shortest? The answer is that we cannot be sure unless we check all other paths. In fact, the shortest path is; definitely not the solution. It is the salesman's problem.

There are two examples. The difference between any-path problems and best-path problems. Best-path problems; are, in general, computationally harder than any-path problems. Any-path problems can be solved in a reasonable amount of time by using heuristics that find good paths to explore. (See the discussion of search in Chapter 3 for one way of doing this.) In the heuristic, are not perfect, the search for a solution may not be a direct path, but that does not matter. For true best-path problems, however, no heuristic that might possibly find the best path can be used. So a much more exhaustive search will be required.

⁷ Of course, representing the cities as a set of coordinates and a recursive procedure could find the shortest path. The algorithm also returns the explicit path of other facts, such as, "city A implies newt alive" or "city B implies not alive".



2.3,5 is the Solution. a State or a Path?

on der the pin/4cm 4.51 finding 'a wit...iment inielpreiatigm ror illie

[lie bank president ale d iii2;11-1311ms,[a !,;.1].H.1 with thlf fork_

There are swctual comporierukti f t hi..ientenee, each which., in isolatium, muy more than Oirie inierprelikliron. Hui Eh: 4...ornporionts It1kl.t form. CCALLient he • 1.111¹-1 :411¹ they constrain e.cich (51.lit's itnerprealiOni_ SOITIC of 1.110 SOLIFCC5 n I u.nthiguiLy in larrrence .are the following:

The ward "bank" male refer either to a financial instinition or to a side of a river, But uni \surd one of they.: may have a president

word "dish" is Ow ubjel..1 oldie "cm." It is pos...Able tia.d. a dish likia_ica(cn. But it is riNore likely that the. pa.A.o. salad in the Wsh was men..

- Pasta is. a salad coimining pasta. Hui ilicre are othoT meanings can he formed from pairs or nourrs._ For example, dog. foirid does not normaby contain flo2s_
- liar phrase "with the fork." could modify several parts or the Nerik..nce. In t Fi i case. it modifies verb "amt,"s But, if The. phra....ic had been 'with vegetables,," then the modification structure would bt different. Arid if' the phrase had been "with her fri.1211cis thn milli:tore. would be different gtill

EliecausL• Or au: in,(12:rac [ion iT]10113gi hC 11 tC rpfetati.Ori of tic 42 tucThis thiN SZIlre.oce, &Mlle Seart h may be required to find a complete Interpretation for the sentence.. .E3ut to solve the problem of finding the interpretatium we need to produce only the interpretation itself. No record of the processing by which the interpretallion found is cio.:4.'ssar je.

Cimtrast with die wattr jug problem. lieu it k sufficient tri. repoti that have solved the prthiein ill that the final state is (2, LPL 1''1ir this. kind Ed' problem, whai. vim really irruist reporE is net the Irina! \$.1,318 the path that we found to that state, 11 itis a statement of a solution to this problem most be ae.querice operations. (Nometinies Called *apethe*) [hat produun the final stale.

Thrise two. examples. uzioral 'angina:go ondors.tindLng wind the water jug problem.. illtr-itrato the difrerkince bawern probiLms whose sulation i' a slate of lthe world and problems whosic solution N a path to a same_ Al one level, this difference can be ignared and problimis Can be funmilated as pries in which only a Auk. k required to be reported.. I i WC do this for problems. such n the water jug. then we inosi redlescribe OLIT Stiites SAD t1 t each tie represents a pa i pall' i 4.1 Sr rLliti on rather than j um a sin k stale the world_ So thi Liu eStiOn

- Solitary, in which the computer is given a problem description and produces an answer with no intermediate communication; and with no demand for an explanation of the reasoning process;
- conversational, in which there is interactive communication. The computer can be asked to provide additional assistance in the problem or to provide a final answer to the user, or both.

Of course, this distinction is not a strict one depending on particular problem domains. As the classic 7-tower problem, mathematical theorem proving could be regarded as either. But for a particular problem, one or the other of these types of systems will usually be decided; indeed, that decision alone is a factor in the choice of a problem-solving method.

2.3.8 Problem Classification

When actual problems are examined from this point of view, several broad classes into which they fall. These classes can each be associated with a generic control strategy that is appropriate for solving the problem. For example, consider the generic problem of *classification*. The task here is to examine an input and decide which of a set of known classes the input is an instance of. Must diagnose, including medical diagnosis, as well as diagnosis of faults in mechanical devices, are examples of classification. Another example of a familiar strategy is *planning*. A design and planning problem can be attacked with this strategy.

Depending on the granularity of the problem, the types of problems and control strategies, many can be grouped with different sets of generic tasks and procedures. See Chandrasekaran (1986) and McDermott (1981) for two approaches to constructing such a hierarchy. It is important to remember here, though, since we are limited to this discussion, that there is no single way of solving all problems. But neither must each new problem be considered totally new. Instead, we analyze our problems carefully and sort our problem-solving methods by the kinds of problems to which they are applicable. We will be able to bring to each new problem much of what we have learned from solving other, similar problems.

2.4 PRODUCTION SYSTEM CHARACTERISTICS

We have just examined a set of characteristics that distinguish various classes of problems. We have also argued that notation is a very good way to describe the operations that can be performed in a search for a solution to a problem. Two questions are inevitably insolubly left at this point are:

1. Can production systems, like problems, be described by a set of characteristic features that stir some light on how they can easily be implemented?
2. If so, what relationship exists between the types of problems and the types of production systems that are suited to solving the problems?

The answer to the first question is yes, for under the following definitions of production systems. A *reducible* production system is one in which the application of a rule never prevents the later application of another rule; that could also have been applied if the first rule was selected. A *monotonic* production system is one in which this is not true. A *partial* production system is a production system with the property that if the application of a particular sequence of rules transforms a state into a final state, then any permutation of those rules that is allowable (i.e., each rule's preconditions are satisfied when it is applied) also maintains the state. A *commutative* production system is a production system that is both monotonic and partially commutative.

³ This corresponds to the definition of a commutative production system in Nilsson (1980).

The significance of these categories or production systems is in the relationship between the categories and their appropriate implementation strategies. But before going into that, it may be helpful to make the meanings of the definitions clearer by showing how they relate to specific problems.

This is the question above, which asks whether there is an interesting relationship between the various classes of production systems and the kinds of problems. For any system, there exist an infinite number of production systems that describe ways to find solutions. Some will be more natural or efficient than others. Any problem that can be solved by a production system can also be solved by a more powerful one. For example, a system that can solve a problem can also solve a more complex problem. One may be surprised to find that a system that is designed to solve a specific problem can also be used to solve other problems.

use individual in represent various sequences or applications of rules of a computer, noncommutative system. So, formal series. there is no relationship between the kinds of problems and the kinds of production systems since all problems can be solved by all kinds of systems. But in a limited sense, there definitely is such a relationship between the kinds of problems and the kinds of systems that lend themselves naturally to describing them. For example, a system that is designed to solve a specific problem can also be used to solve other problems. This is the case for the four categories of production systems: monotonic, nonmonotonic, partial, and commutative.

	Monotonic	Nonmonotonic	Partial	Commutative
commutative	it is proving.	Robot navigation		
Not commutative	Chemical synthesis	Bridge		

commutative versus.

Fig. 2.17 The Four Categories of Production Systems.

nonpartially commutative, along with some problems that can be solved by each type of system. The upper left corner represents commutative systems.

Partially commutative, problems that are solvable by a partially commutative system. This is a surprising result. For example, a problem that is solvable by a partially commutative system is also solvable by a noncommutative system. But recall that ignorable problems are those for which a natural formulation leads to a solution that can be ignored. Such a natural formulation will then be a partially commutative, monotonic system. Problems that involve creating new things rather than changing old ones are generally solvable. We have described it, in one example, as a creative process. But those processes can also be implemented with a partially commutative, monotonic system.

Partially commutative, monotonic production systems are important from an implementation standpoint because they can be implemented without the ability to backtrack. It is important to note that an incorrect path has been followed. Although it is possible to implement such systems with backtracking in order to guarantee the solution, this often results in a considerable increase in efficiency. Particularly for a database since the database will never have to be restored. It is not necessary to keep track of where in the process every change was made.

Well, the main reason for this is that the ability to backtrack is not necessary. For example, a robot that is moving in a straight line can be controlled by a partially commutative, monotonic system. But recall that ignorable problems are those for which a natural formulation leads to a solution that can be ignored. Such a natural formulation will then be a partially commutative, monotonic system. Problems that involve creating new things rather than changing old ones are generally solvable. We have described it, in one example, as a creative process. But those processes can also be implemented with a partially commutative, monotonic system.

on how the op.12riatorsuri: choscri. thi8Plizzli2 arid thL. blocks world problem can also be considered partially commutative.

Both types of partially commutative production systems are significant from an implementation point of view because they tend to lead to many duplications of individual states, during the generation process. This is discussed further in Section 2.5.

Production systems that are not partially commutative are useful for many problems in which irreversible changes occur. For example, consider the problem of determining a process to produce a desired (chemical) compound. The operators available include such things as 'Add acid to the pot' or 'Orange the temperature to 1 degrees.' The operators may cause irreversible changes to the solution being brewed. The order in which they are performed can be very important in the final output. It is possible that if x is added to y , a stable compound will be formed, so later addition of z will have no effect; if z is added to y , however, a different stable compound may be formed, so later addition of x will have no effect. Noncommutative production systems are less likely to produce the same node many times in the search process. When dealing with ones that include irreversible processes, it is important to make corrections

the ii Est ii ri lc , although in the initial stages, planning can help to make these less important.

2.5 ISSUES IN THE DESIGN OF SEARCH PROGRAMS

Every search process can be viewed as a traversal of a tree in which each node represents a problem state, and each arc represents a transition between states. For example, Fig. 2.18 shows part of a search tree for a water pouring problem. The arcs have not been labeled in the Fig., but they correspond to particular water-pouring operations. The search process must find a path or paths through the tree that connect an initial state with one or more final states. The tree that must be searched could, in principle, be very large, from the rules of the problem. But, in practice, most of it is never explored.

of it is never explored. The search process builds the tree explicitly and then searches it. Most search processes represent the tree implicitly in the rules and generate it explicitly only as needed. In our discussion of search methods, it is important to keep in mind this distinction between implicit and explicit search.




Fig. 2.1a A Search Tree for the Water Jugs Problem

In the next chapter, we present a family of general-purpose search techniques. But before doing so we need to mention some important issues that arise in search:

- * The direction in which to conduct the search (forward or backward). We can search forward through the state space from the start state to a goal, or we can search backward from the goal.
- The width of the search. The search process should not spend too much of its time looking

for nodes to apply, so it is critical to have efficient procedures for matching rules against states.

- The representation of the state. A node in the search process (the knowledge represented by the node and the facts in the world) can be represented by a simple structure, or it can be a complex structure. For problems like chess, a node can be represented by a simple structure, or it can be a complex structure. For problems like chess, a node can be represented by a simple structure, or it can be a complex structure. For problems like chess, a node can be represented by a simple structure, or it can be a complex structure.

We discuss the knowledge representation and time problems further in Chapter 4. We investigate matching and forward versus backward reasoning where we return to production systems in Chapter 6.

One other issue we should consider at this point is that of search trees versus search graphs. As mentioned above, we can think of production rules as generating nodes in a search tree. Each node can be expanded in turn, generating a set of successors. This process continues until a node representing a solution is found. Implementing such a procedure requires little bookkeeping. However, this process often results in the same node being generated several times over several paths and so being processed more than once. This happens because the search space may really be an arbitrary directed graph rather than a tree.

For example, in the tree shown in Fig. 2.18, the node (4,3), representing 4-gallons of water in the jug and 3 gallons in the other, can be generated either by first filling the 4-gallon jug and then the 3-gallon one or by filling the 3-gallon one first and then the 4-gallon one. The order does not matter. Continuing to process both these nodes would be redundant. This example also illustrates another problem that often arises when the search process operates as a tree walk. On the third level, the node (0,0) appears (In fact, it appears twice.) But this is the same as the root node of the tree, which has already been expanded.

Those two paths have not gotten us anywhere. So we would like to eliminate them and continue only along the other branches.

Instead of traversing a search tree, we traverse a

The waste of effort that arises when the same node is generated more than once can be avoided at the expense of additional

differs from a tree in that several paths

may come together at a node. The graph corresponding to the Fig. 2.1) A search graph is the wave-

tree of Fig. 2.1 is shown in Fig. 2.19.

Any tree search procedure that keeps track of all the nodes that have been generated so far can be converted to a graph search procedure by modifying the action performed each time a node is generated. Notice that of the two systematic search procedures we have discussed so far, this requirement that nodes be kept track of is met by breadth-first search but not by depth-first search. But, of course, depth-first search could be modified, at the expense of additional storage, to retain in memory nodes that have been expanded and then backed-up over. Since all nodes mentioned in the search graph, we must discuss the following algorithm instead of simply adding a new node to the graph,

Algorithm to Check Duplicate Nodes

1. Examine the set of nodes that have been created so far to see if the new node already exists.
2. If it does not, simply add it to the graph just as for a tree.
3. If it does already exist, then do the following:
 - (a) Set the node that is being expanded to point to the already existing node corresponding to its successor rather than to the new one. The new one can simply be thrown away.
 - (b) If you are keeping track of the best (shortest or otherwise least-cost) path to each node, then check to see if the new path is better than the old one. If not, do nothing. If better, record the new path as the correct path to the node and propagate the corresponding change in cost down through successor nodes as necessary.

One problem that may arise here is that cycles may be introduced into the search graph. A cycle is a path through the graph in which a given node appears more than once. For example, the graph of Fig. 2.19 contains two cycles of length two. One includes the nodes (0,0) and (4,0); the other includes the nodes (0,0) and (0,3). Whenever there is a cycle, there can be paths of arbitrary length. Thus it may become more difficult to know that a graph traversal algorithm is guaranteed to terminate.

Treating the search process as a graph search rather than as a tree search reduces the amount of effort that is spent exploring essentially the same path over and over. But it incurs additional effort each time a node is

generated to see if it has beer; generated berure. 'Whether this effort is justified depends on particular problem. If it is likely that the same node will be generated in several different ways, then it is more worthy: hire IR.) use a graph procedure than if such duplication will happen only rarely.

Graph search procedures are especially useful for dealing with partial, non-commutative production system in which a set of operations will produce the same result regardless of the order in which the operations are applied. A systematic search procedure will try many of the permutations of these operators and so will generate the same node many times. This is especially true for the *jug* example shown at the end.

24 ADDITIONAL PROBLEMS

Several problems have been discussed throughout this chapter. Other problems have not yet been mentioned, but are common throughout the AI literature. Some have been in classic AI books, but could be easily added to this section. A useful exercise at this point, would be to visit each of them in turn and try to solve them. We have 'jug' discussed.

A heuristic (List) is a problem which is rare this parade of toy problems is presented. Artificial Intelligence is not nearly as simple as in problems and microworlds (such as blocks world) or the (unique) them have been used for a long time, so problems have been the core of systems which solve very novel problems. So think about these problems not as defining the scope of AI but rather as providing a core from which much can be developed.

The Missionaries and Cannibals Problem.

Three missionaries and three cannibals find themselves on one side of a river. They have agreed that they would all like to get to the other side. But the missionaries are afraid that the cannibals have agreed to. So the missionaries make the trip across the river in Nadi alai. The number of missionaries on the same side of the river is never less than the number of cannibals who are on the same side. The only boat available has only two people at a time. It can never cross the river without the missionaries, risking being eaten?

The Tower of Hanoi

There are three towers near Hanoi where the monks live. They devote their lives to a very simple task. In their monastery, there are three towers. On these posts is a stack of four disks, each with a hole in the center. The disks are of different sizes, the largest at the bottom and the smallest at the top. The monks' task is to move all of the disks to one or the other pegs. Only one disk can be moved at a time, and all the other disks must be on one of the other pegs. Only one disk can be placed on top of a smaller disk. The third peg can be used as a temporary resting place for the disks. What is the shortest sequence of moves to accomplish their task? Even the best solution to this problem will take the monks a long time. This is fortunate. Nine legend has it that the world will end when they have

The Monkey and Bananas Problem

A hungry monkey finds himself in a room in which a bunch of bananas is hanging from the ceiling. The monkey is 1.3 m tall, can reach 1 m, and the bananas are 2.3 m high. In the room there are a chair and a stick. The ceiling is just the right height so that a monkey sitting on a chair could knock the bananas down with the stick. The monkey knows how to move around, carry other things around, reach for the bananas, and wave a stick in the air. What is the shortest sequence of actions for the monkey to take to acquire lunch?

SEND	DONALD	CROSS
+MORE	+GERALD	+R ADDS
L A L L A L 1 1 1 1 1 1		
MONTY	ROBERT	DANGER

Fig. 2.20 Some Cryptarithmic Problems

Cryptarithmic method

Consider an arithmetical problem represented as shown in the examples in Fig. 2.20. Assign a decimal digit to each of the letters in such a way that the answer to the problem is correct. If the same letter occurs more than once, it must be assigned the same digit each time. No two different letters may be assigned the same digit.

Several strategies for solving cryptarithmic problems have been, studied intensively by Newell and Simon (1972).

SUMMARY

In this chapter, we have discussed the first two steps that must be taken toward the design of a program to solve a particular problem!

1. Define the problem precisely. Specify the problem space, the operators for moving within the space, and the starting and goal state(s).
2. Analyze the problem to determine where it falls with respect to seven important issues,

The last two steps for designing a program to solve that problem are, of course:

3. Identify and represent the knowledge required by the task.
4. Choose one or more techniques for problem solving, and apply those techniques to the problem.

Several general-purpose techniques are presented in the next chapter, and several of them have already been alluded to in the discussion (if the problem characteristics in this chapter. The relationships between problem characteristics and specific techniques should become even clearer as we go on. There, in Part II, we discuss the issue of how domain knowledge is to be represented.

EXERCISES

In this chapter, the following problems were mentioned:

- Chess
- 8-puzzle
- Missionaries and cannibals
- Monkey and bananas
- Water jug
- Traveling salesman
- Towers of Hanoi
- Cryptarithmic

Analyze each of them with respect to the seven problem characteristics discussed in Section 2.3.

2. Before we can solve a problem using state space search, we must define an appropriate state space. For each of the problems mentioned above for which it was not done in the text, find a good state space representation.
3. Describe how the branch-and-bound technique may be used to find the shortest solution to the Water Jug Problem.

4. For each of the following types of problems, try to describe at food he function'.
 - (a) world
 - (b) Theorem prkwing..
 - (c) Missiorrarics. n d cannibals
5. !Give an example of . problem far which hreachill-first search would veni 1.: better thin depth-firm search_

Give an example of a prohle n r ror which tic earch would work: boiler (ham breadth-first search_
- . Write an i4orilhorl to perform breadth-first :44:.-iarch of a problem xraph. Maki: sure your algorithm works properly MI hen a single node i generated a[more than one level in the graph.
7. Try to construct an .r4.oritlim for solving blocks. world prob]rns., such the in in Fig- 2_10_ F ri nnt cheat by tool.....ing ahead Chapter 13,